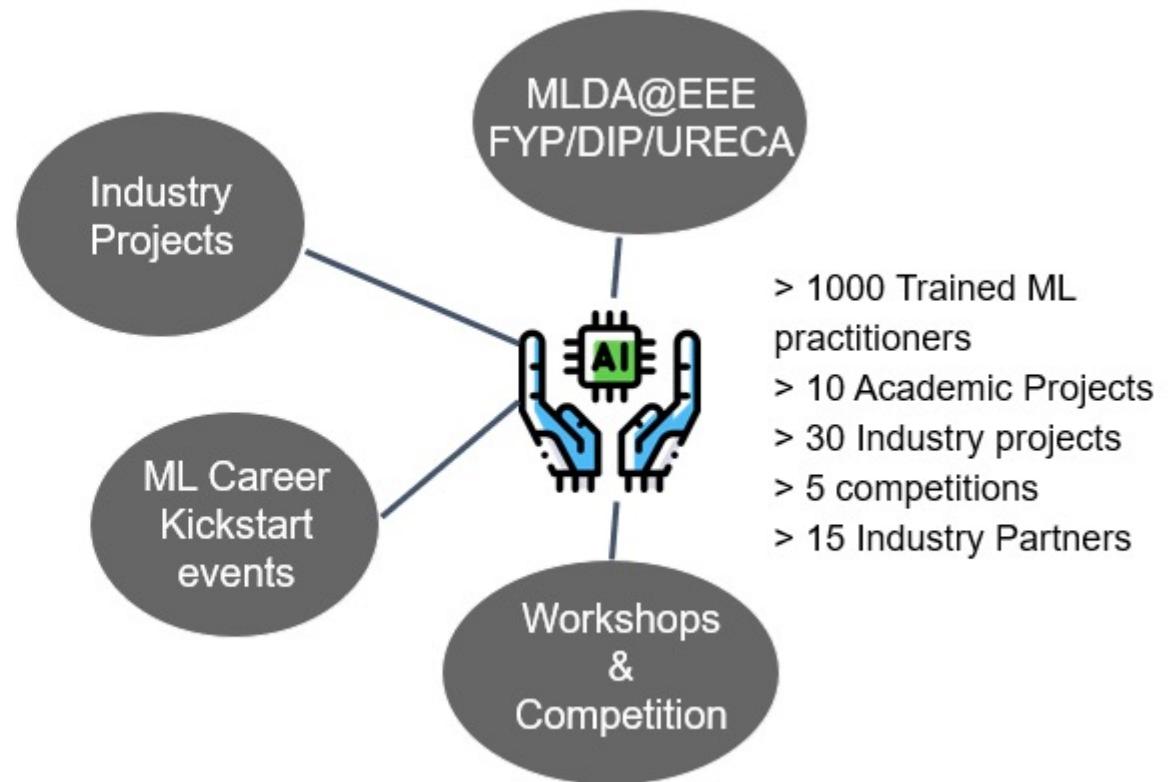
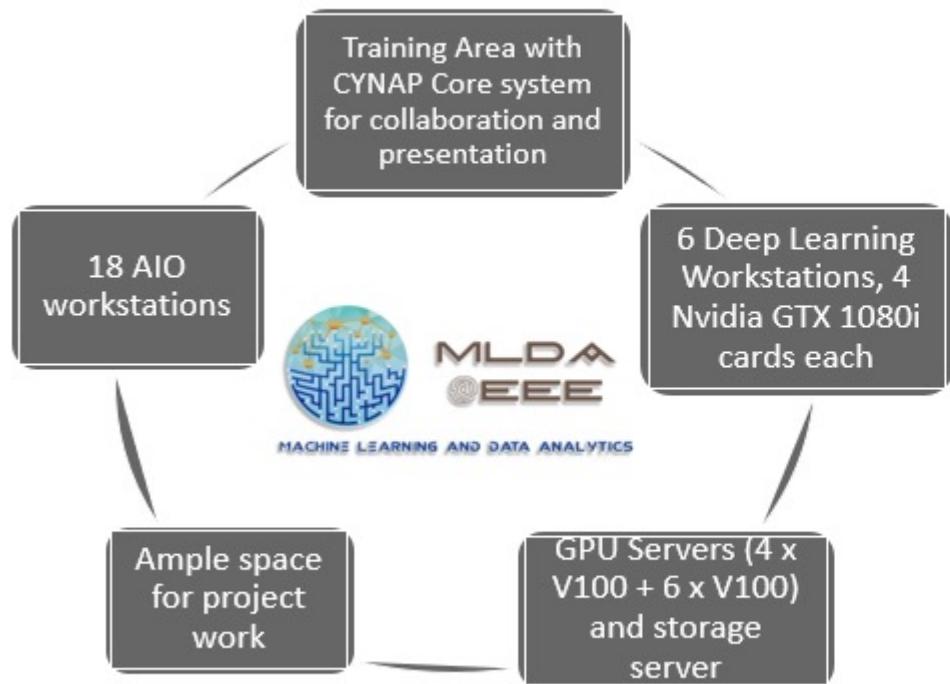


Generative Adversarial Networks (GANs) Workshop

**Presented by:
Dewi & Van**

MLDA's Mission

Provide an integrated platform for EEE/IEM students to learn and implement Machine Learning, Data Science & AI, as well as facilitate connections with the industry.



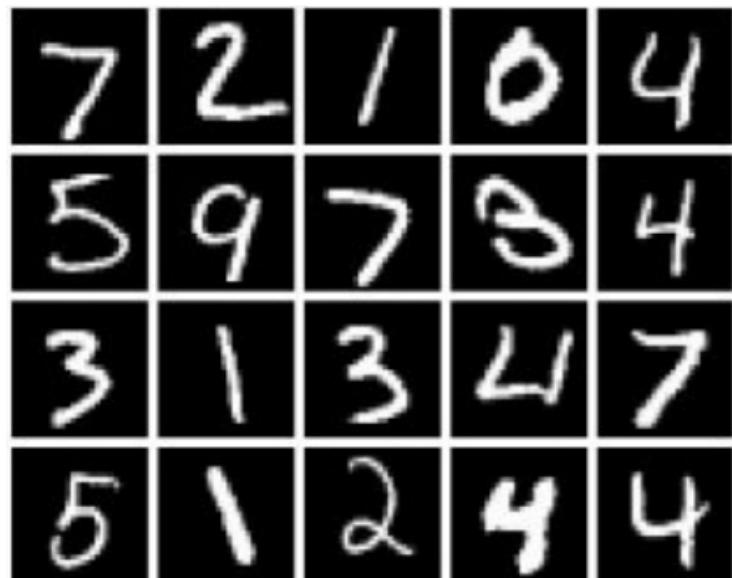
Which one is FAKE?



GENERATIVE ADVERSARIAL NETWORKS (GANs)

Generative Models

Given training data, generative models aim at learning the true data distribution of the training set to generate new data points from this distribution with some variations.



Training samples

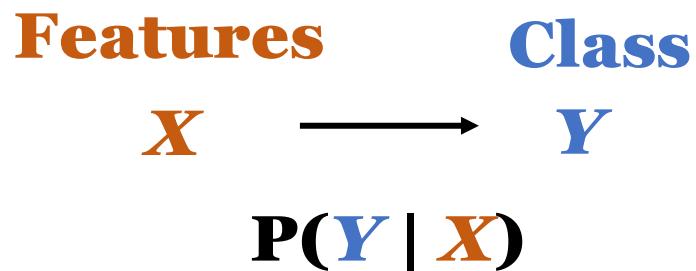
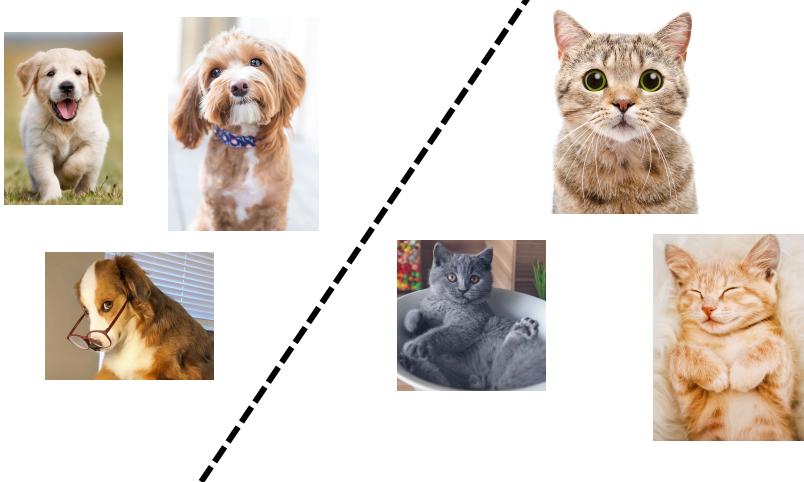


Generated samples

Classifier Model vs. Generative Model

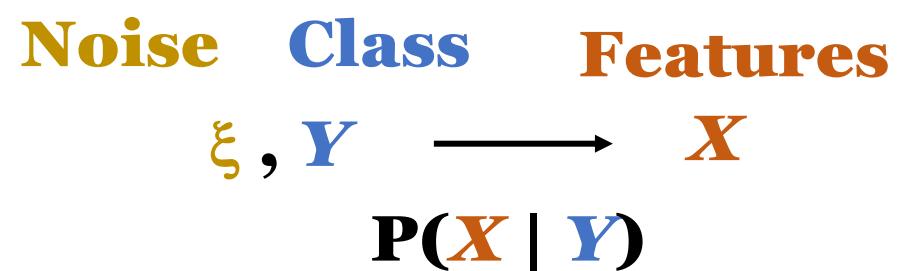
Classifier Model

- Goal: Models the **decision boundary** between classes



Generative Model

- Goal: Models the **actual distribution** of each classes



Generative Models

Given training data, generative models aim at learning the true data distribution of the training set to generate new data points from this distribution with some variations.

Two famous deep generative model algorithms:

- **Variational Autoencoder (VAE)**
- **Generative Adversarial Network (GAN)**

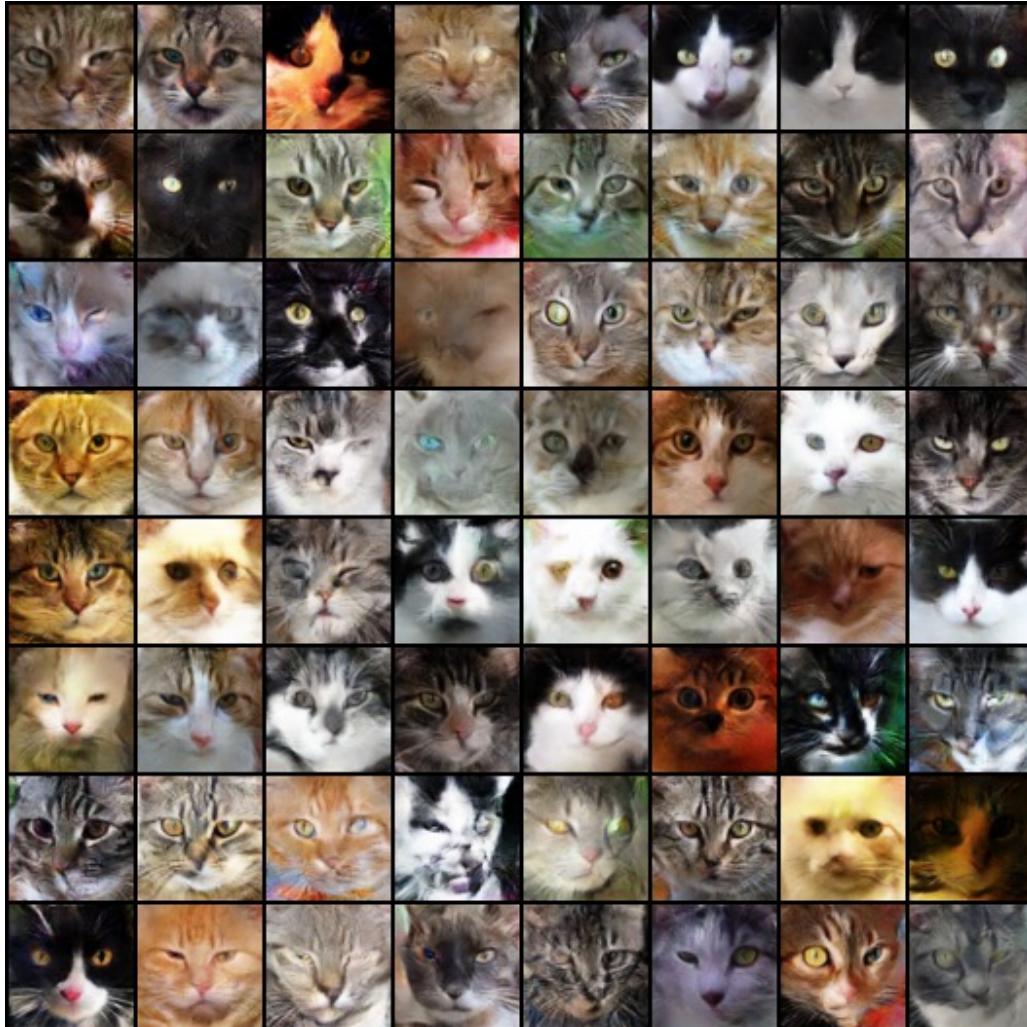
GANs: Cool Applications



**Face Generation
StyleGAN2**

More from
[https://thispersondoesno
texist.com/](https://thispersondoesnotexist.com/)

GANs: Cool Applications



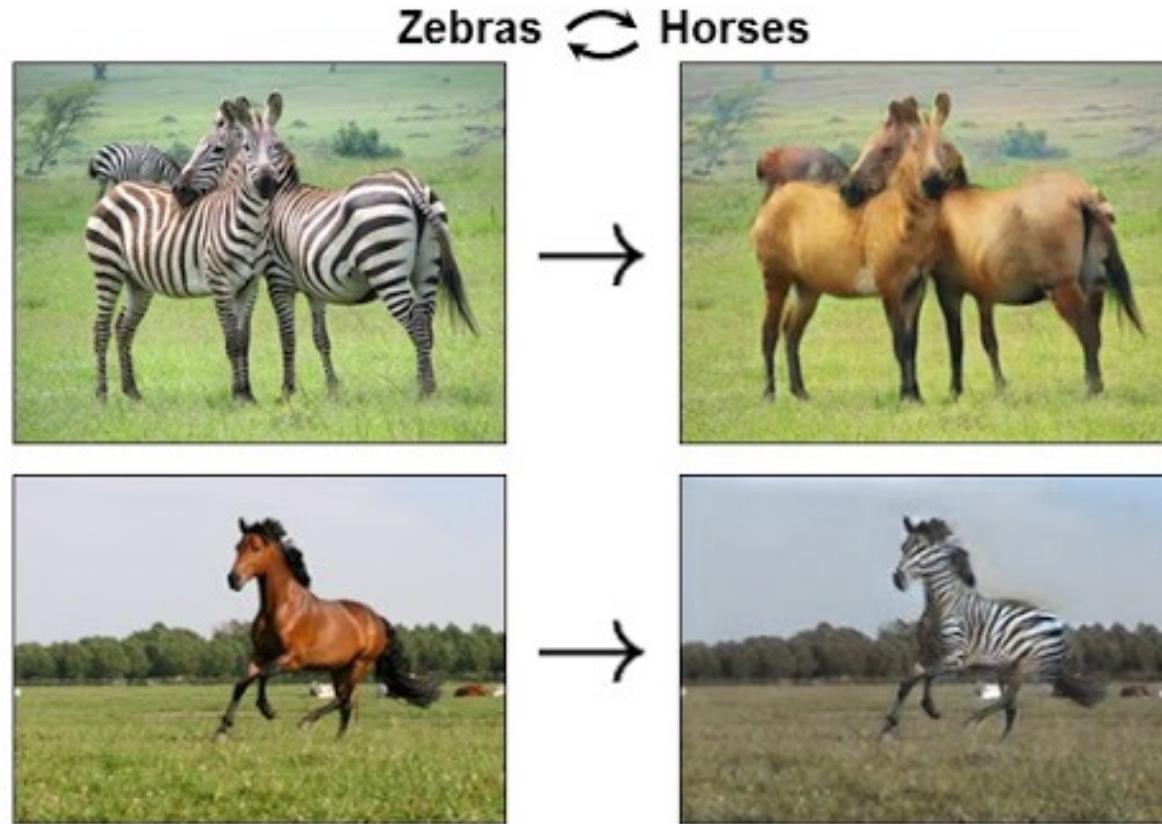
**Cat Generation
StyleGAN2**

GANs: Cool Applications



**Super Resolution
SRGAN**

GANs: Cool Applications



**Image Translation
CycleGAN**

Generative Adversarial Networks

Generator

Learns to make **fakes** that
look **real**

Discriminator

Learns to distinguish **real**
from fake

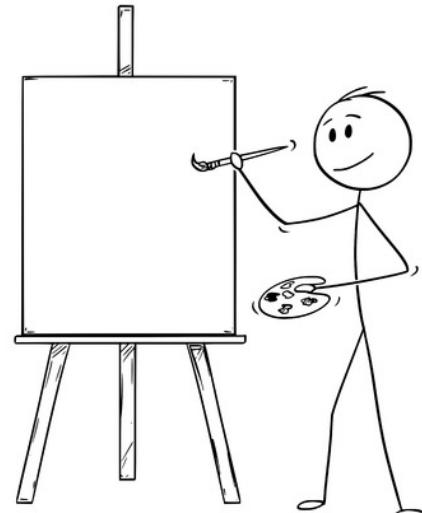
Generative Adversarial Networks

Generator

vs.

Discriminator

Learns to make **fakes** that
look **real**



**Art
Forger**

Learns to distinguish **real**
from **fake**

vs.

**Art
Critic**

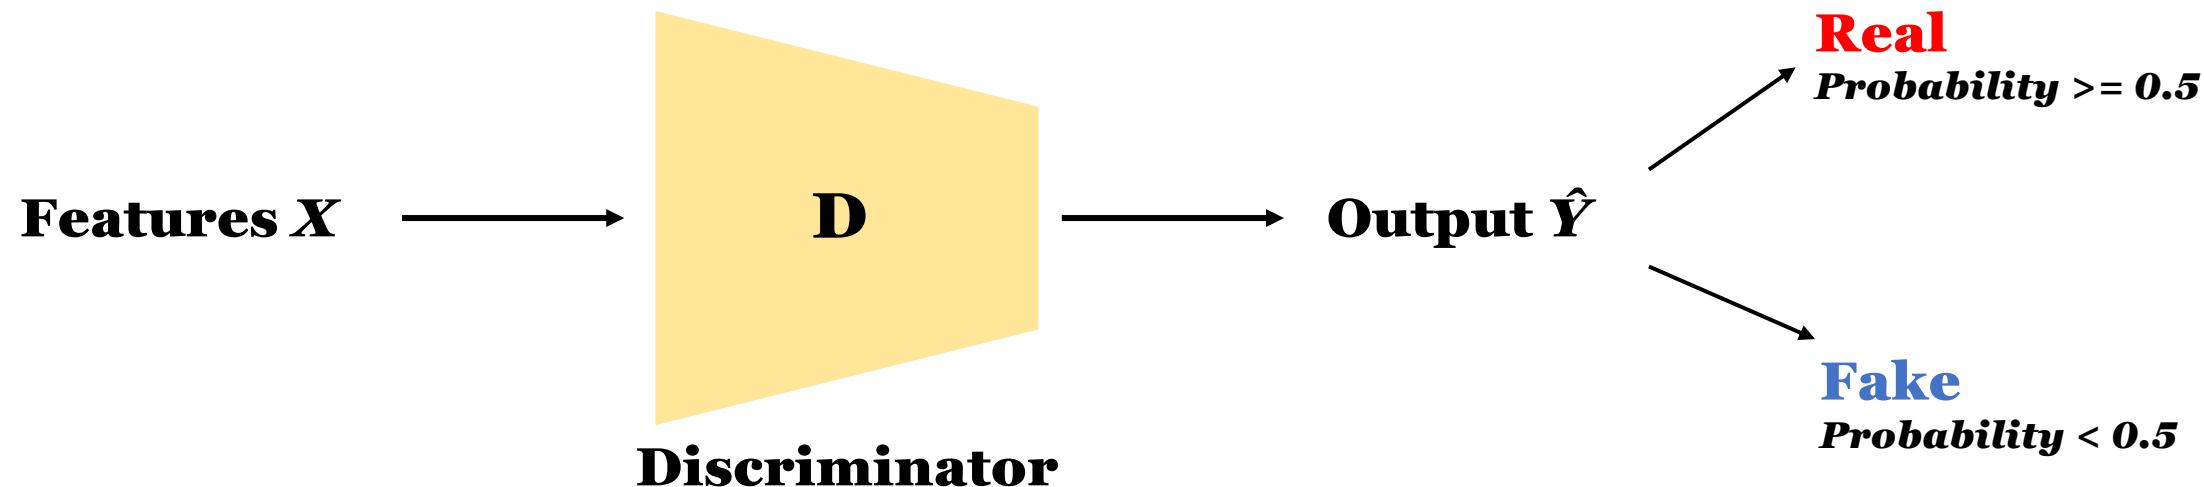


Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

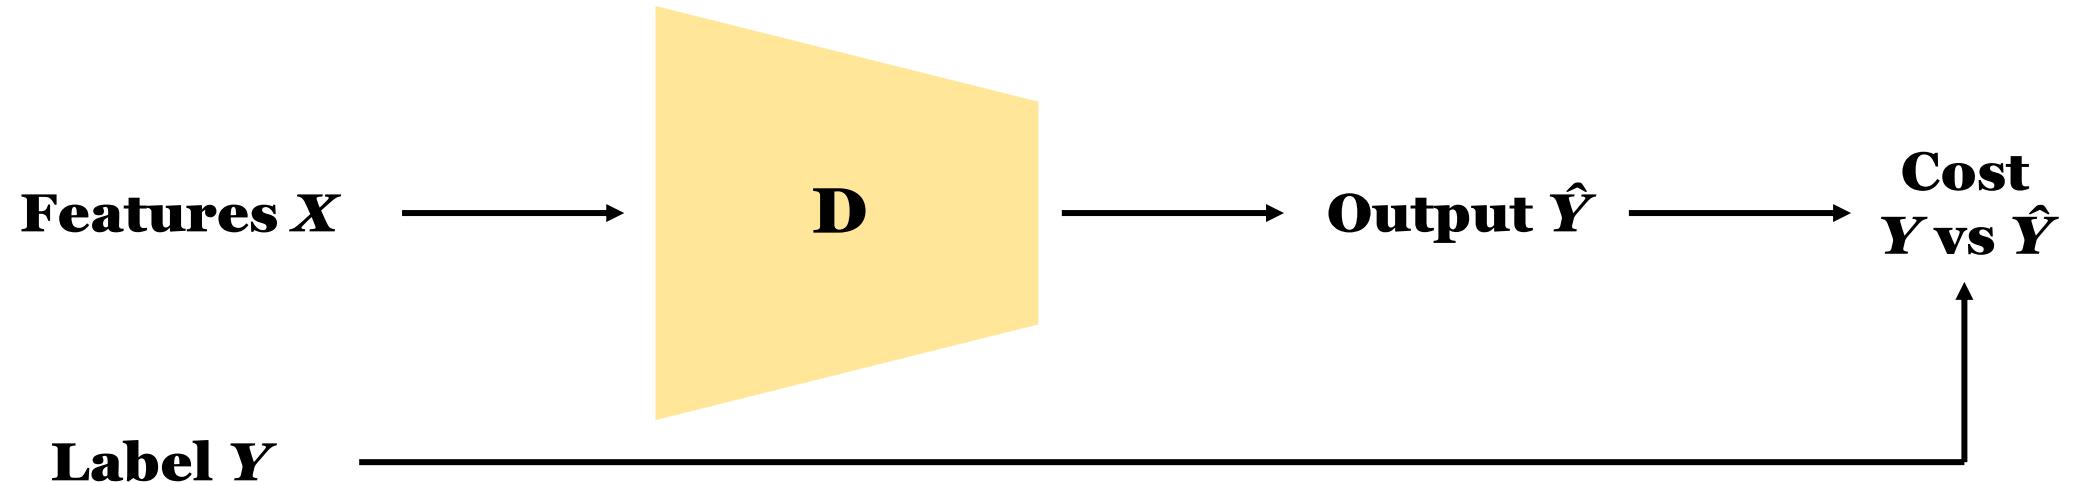


Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

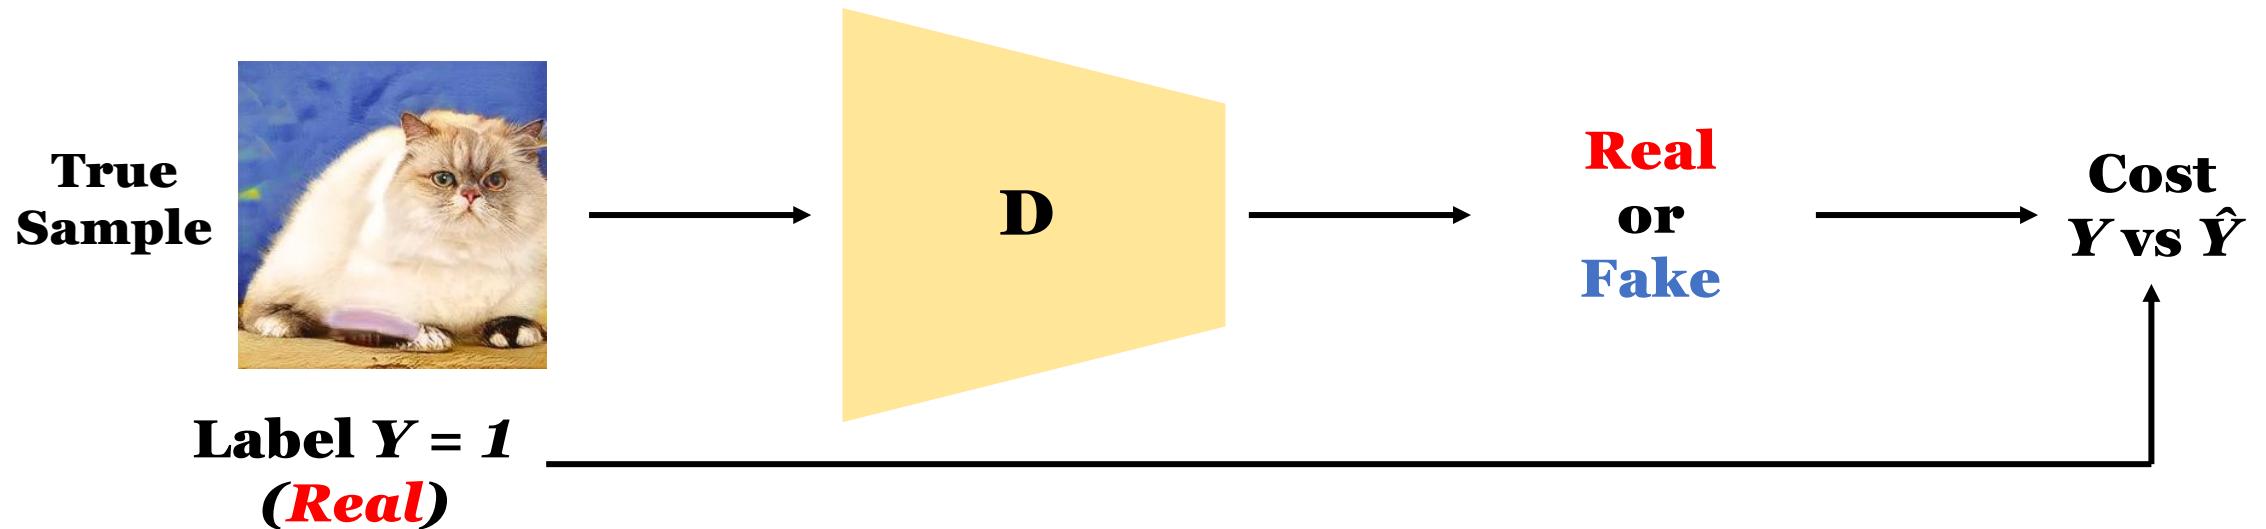


Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

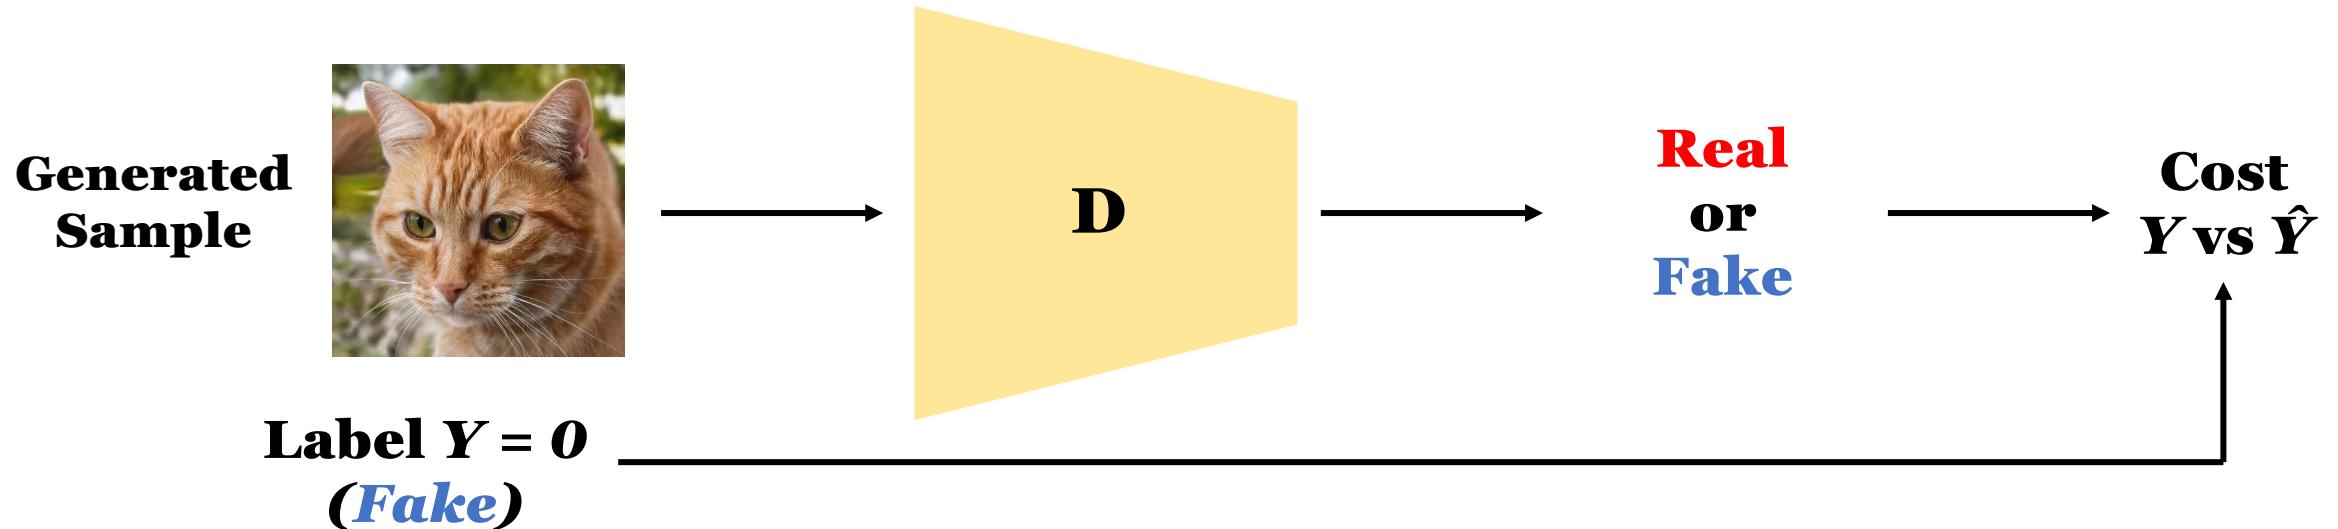


Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

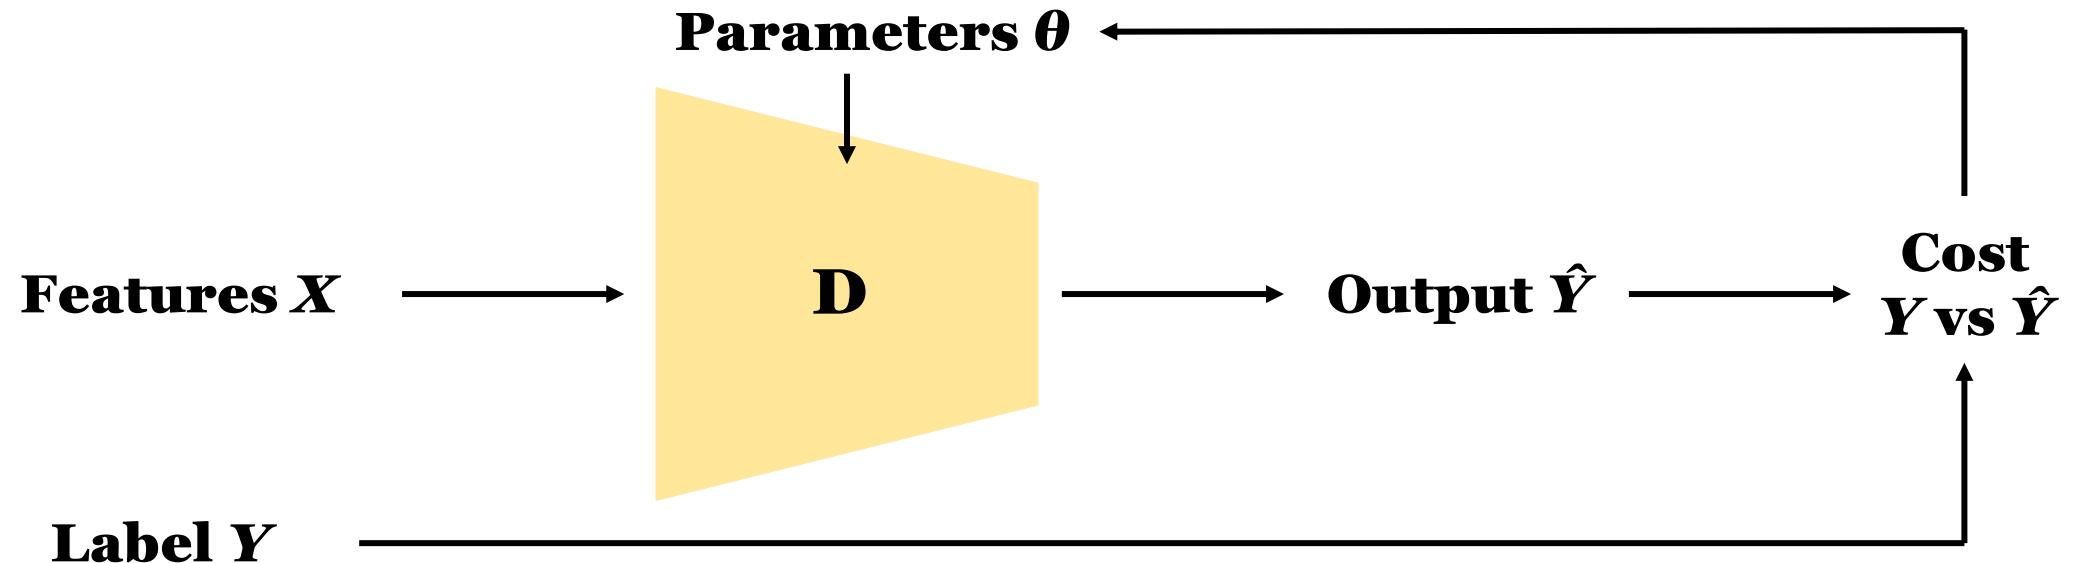


Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real



Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

$$P(Y | X)$$

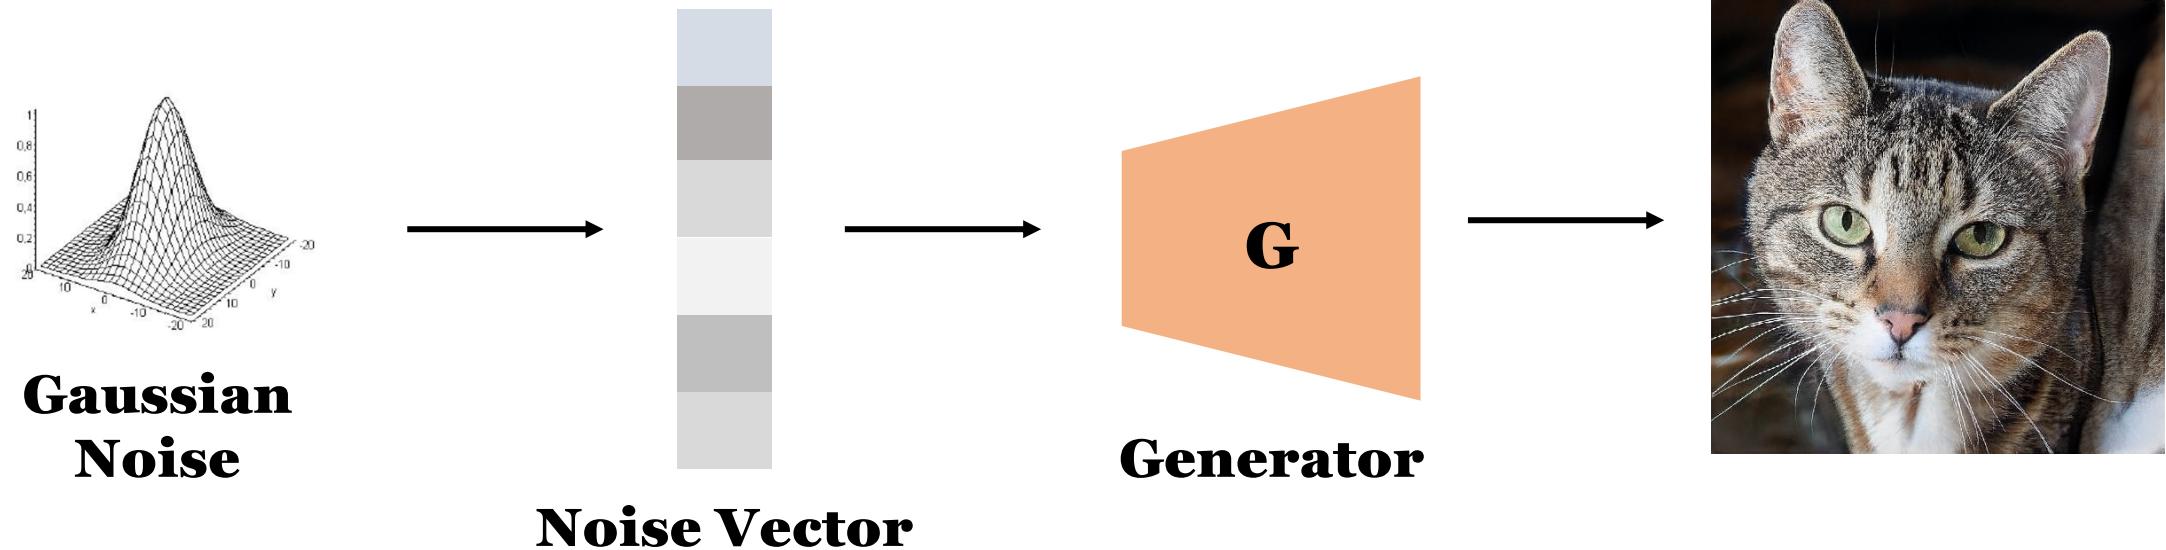
Class Features

Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

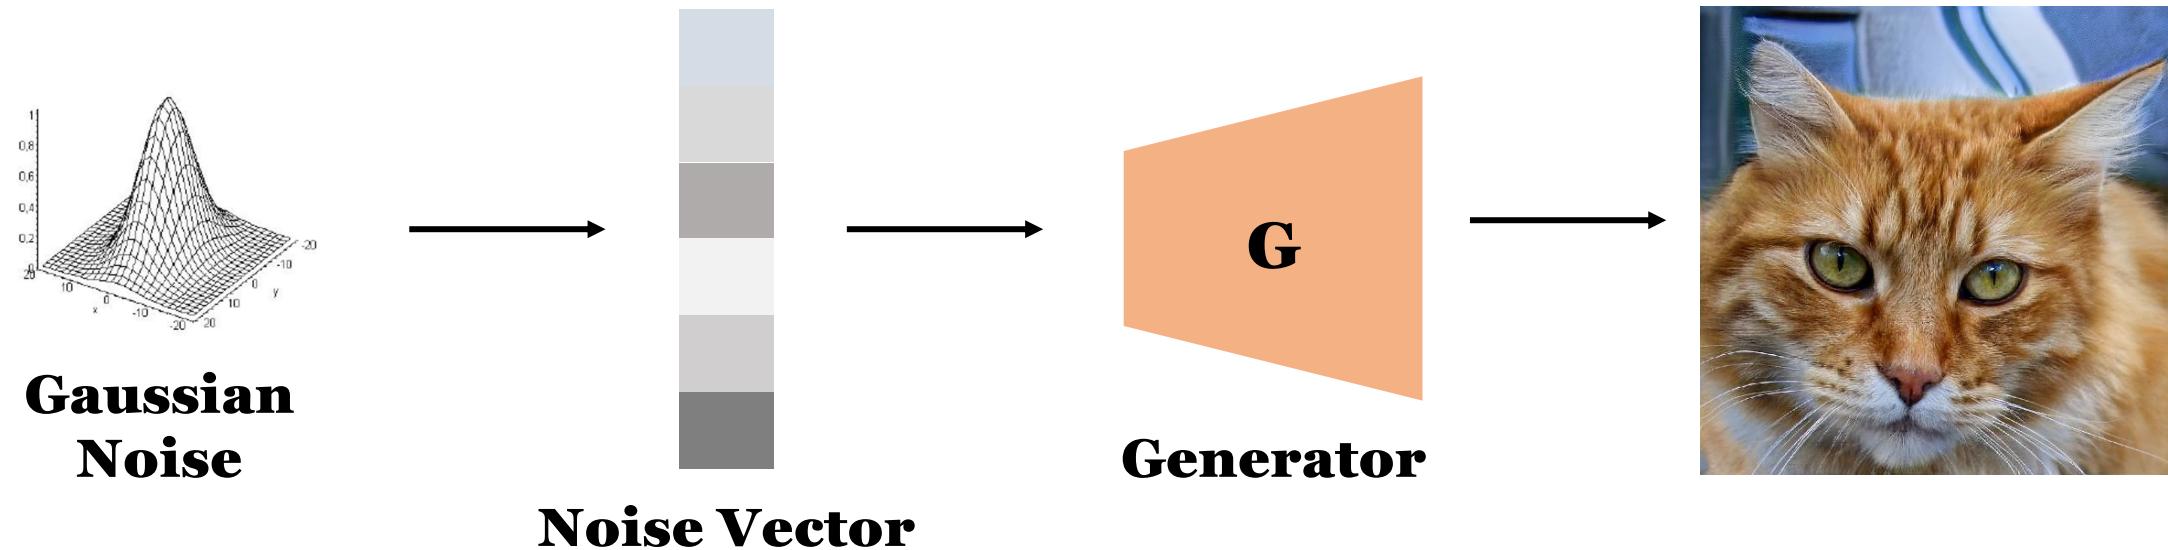


Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Noise Vector

- Also known as **latent variable**, they are those variables that are important for a domain but are not directly observable.
- Random noise vector's dimensionality is smaller than the target output's dimensionality
- Noise helps GAN to produce a wide variety of data

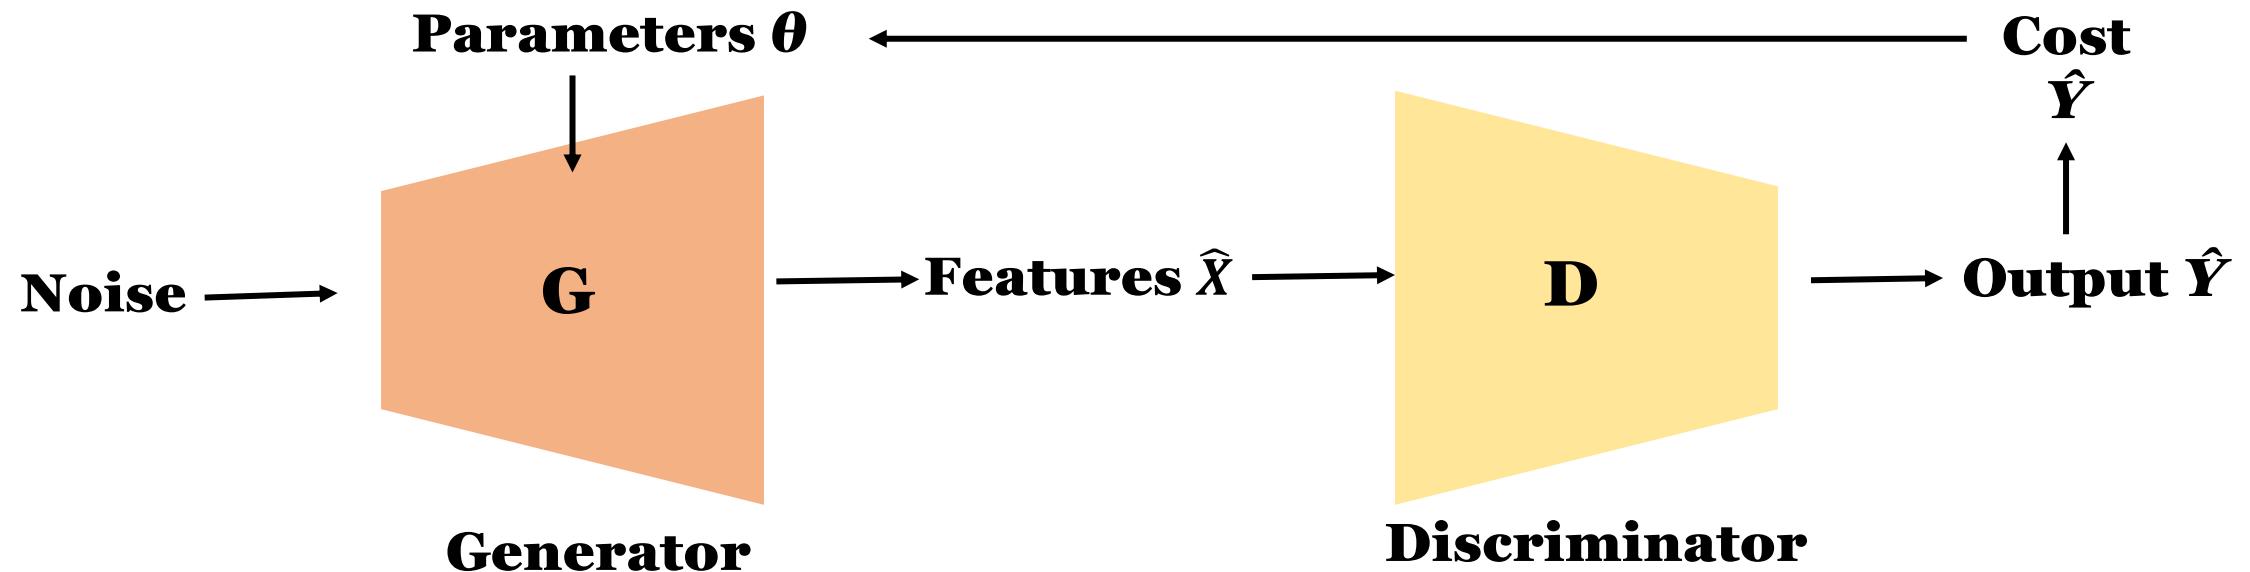
Noise Vector

Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

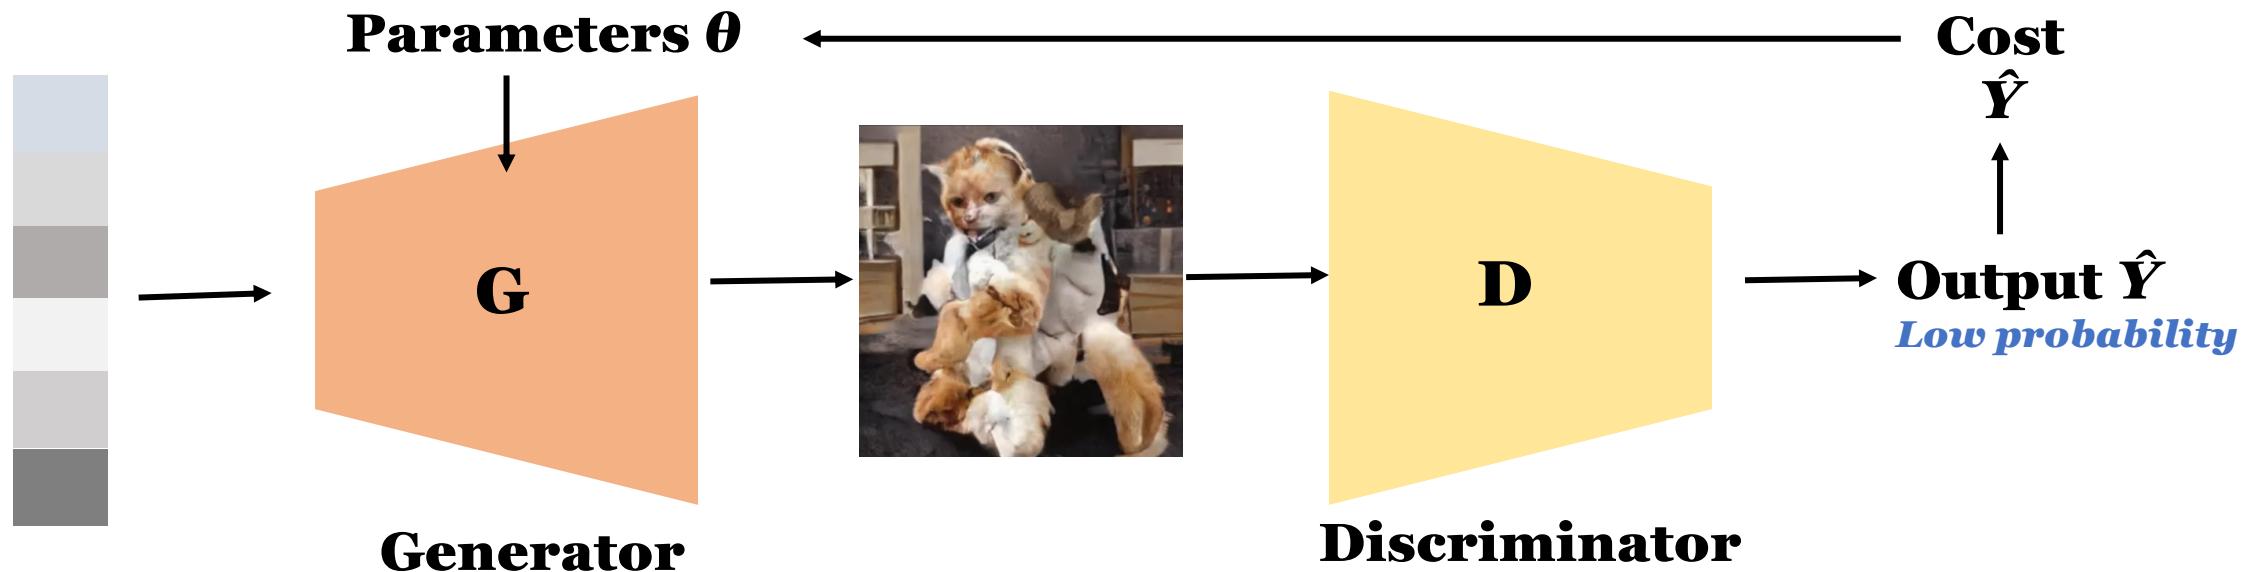


Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

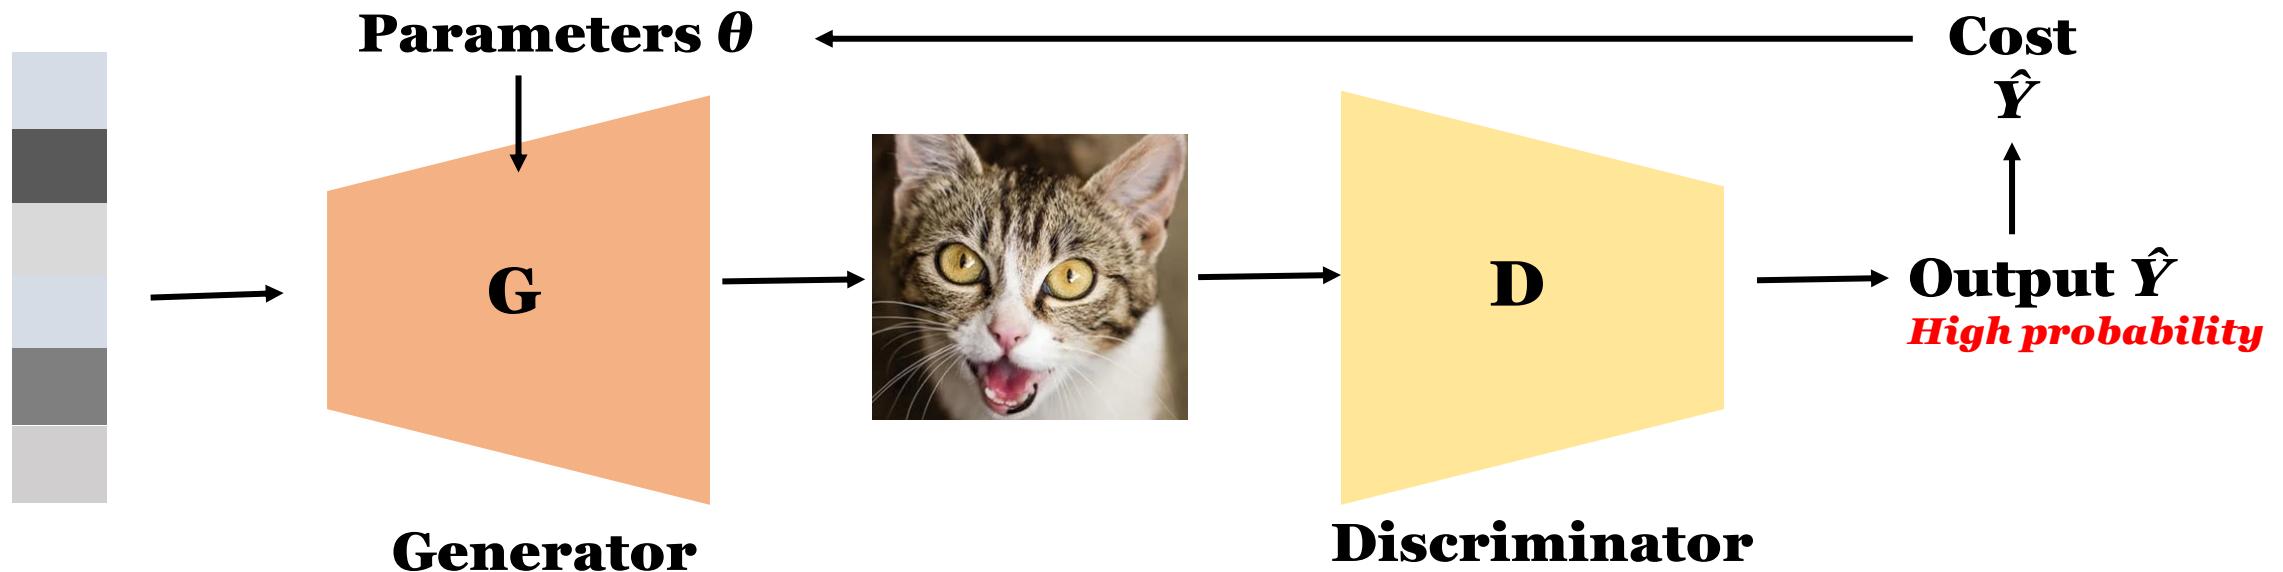


Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

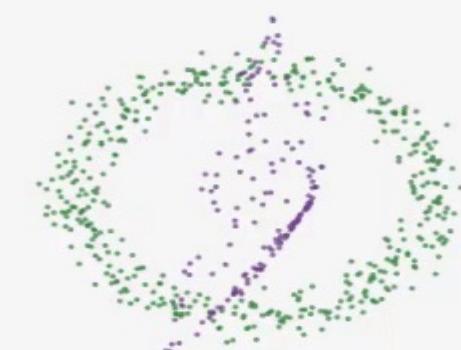
$$P(X \mid Y)$$

Features Class

Approximate the distribution of problem domain



LAYERED DISTRIBUTIONS



Each dot is a 2D data sample:
• Real samples
• Fake samples

<https://poloclub.github.io/ganlab/>

Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

m: sample size
y: label
 \hat{y} : prediction

Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

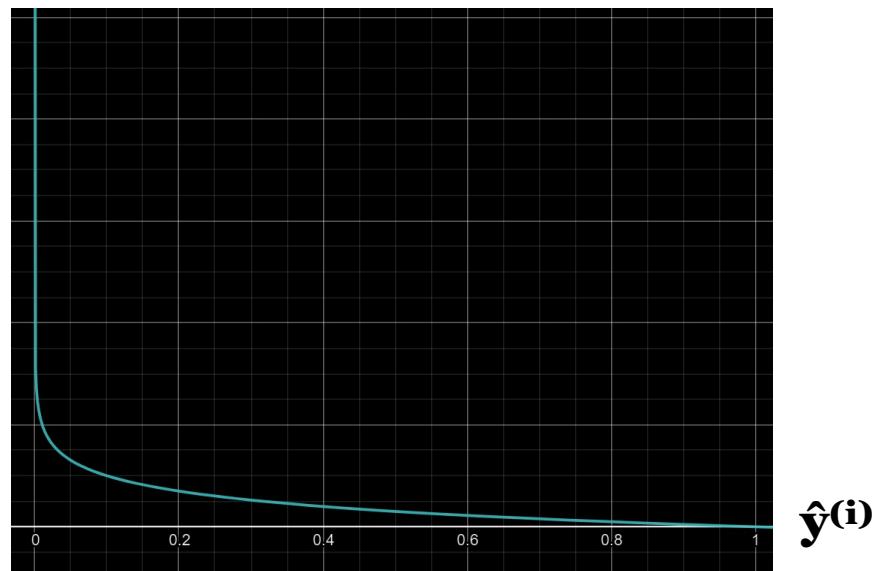
$$BCE = -\frac{1}{m} \sum_{i=1}^m \boxed{y^{(i)} \log \hat{y}^{(i)}} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Label y = 1 (real)

m: sample size
y: label
 \hat{y} : prediction

Note: $\log(1) = 0$

$\log \hat{y}^{(i)}$



Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

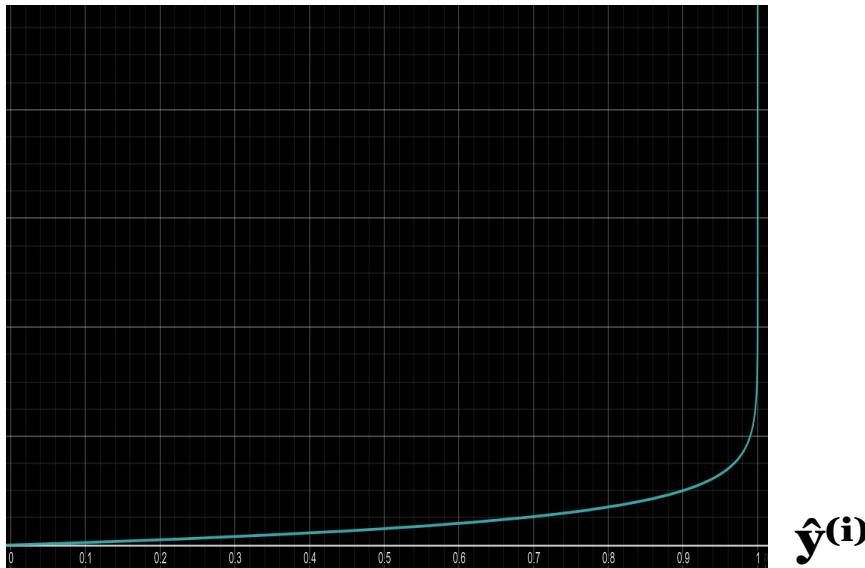
$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + \boxed{(1 - y^{(i)}) \log(1 - \hat{y}^{(i)})}$$

Label $y = 0$ (fake)

m: sample size
y: label
 \hat{y} : prediction

Note: $\log(1) = 0$

$\log(1 - \hat{y}^{(i)})$



Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

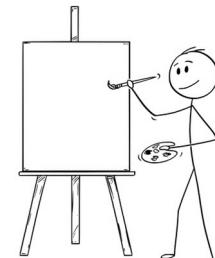
$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

BCE Loss -> Minimax Loss

$$\min_d \max_g - \left[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z)))) \right]$$

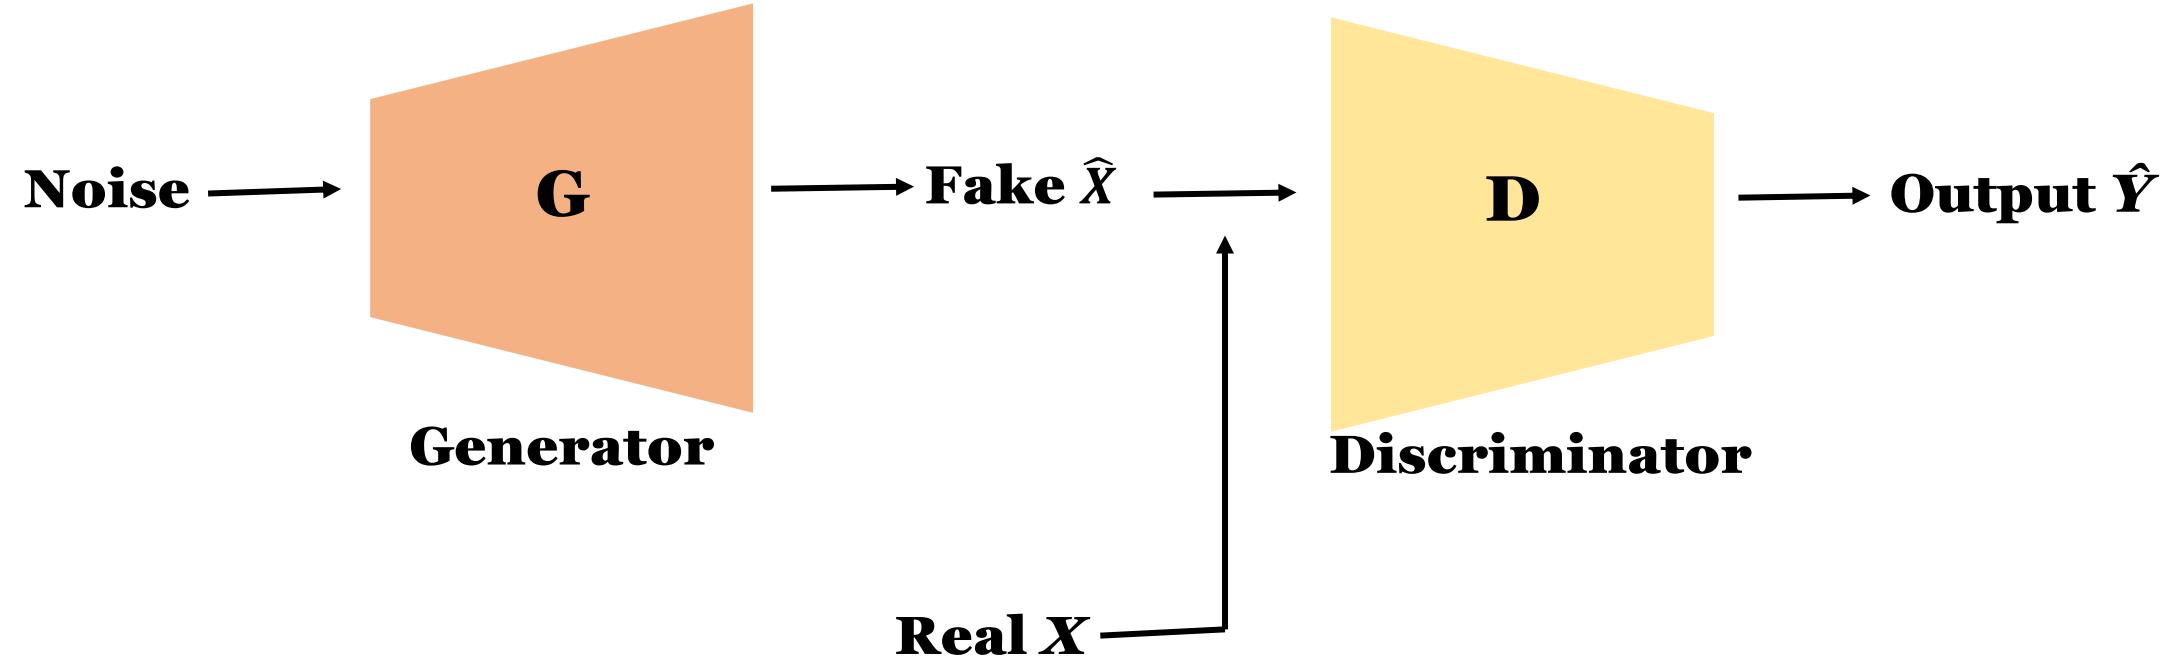


Discriminator:
Minimize cost



Generator:
Maximize cost

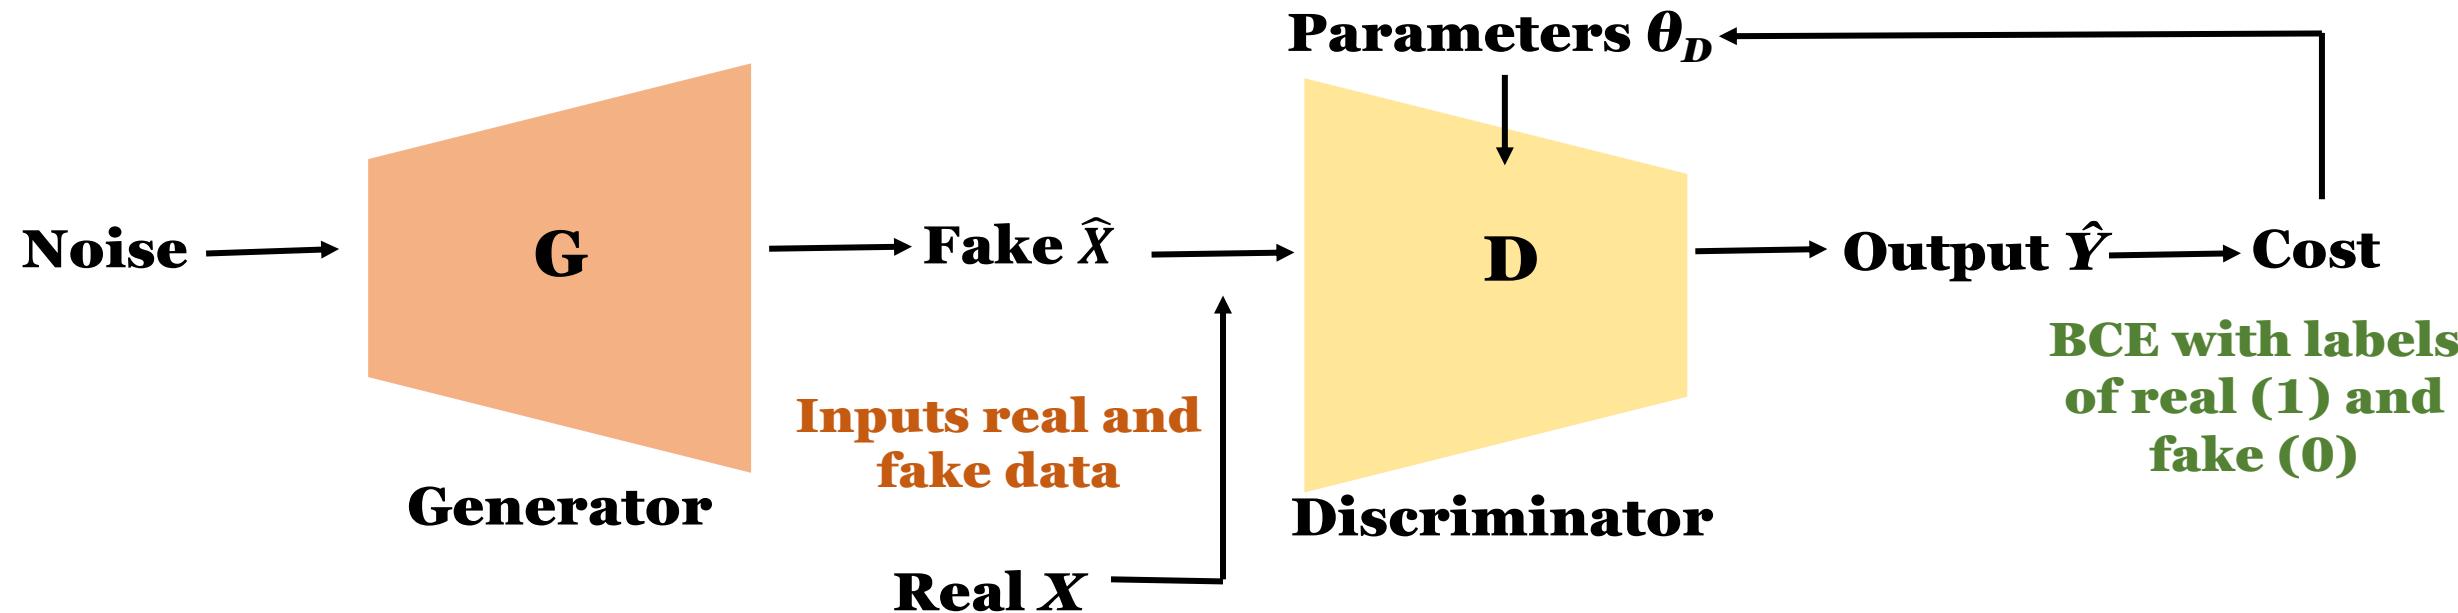
Complete GAN Model



Complete GAN Model

Training Discriminator

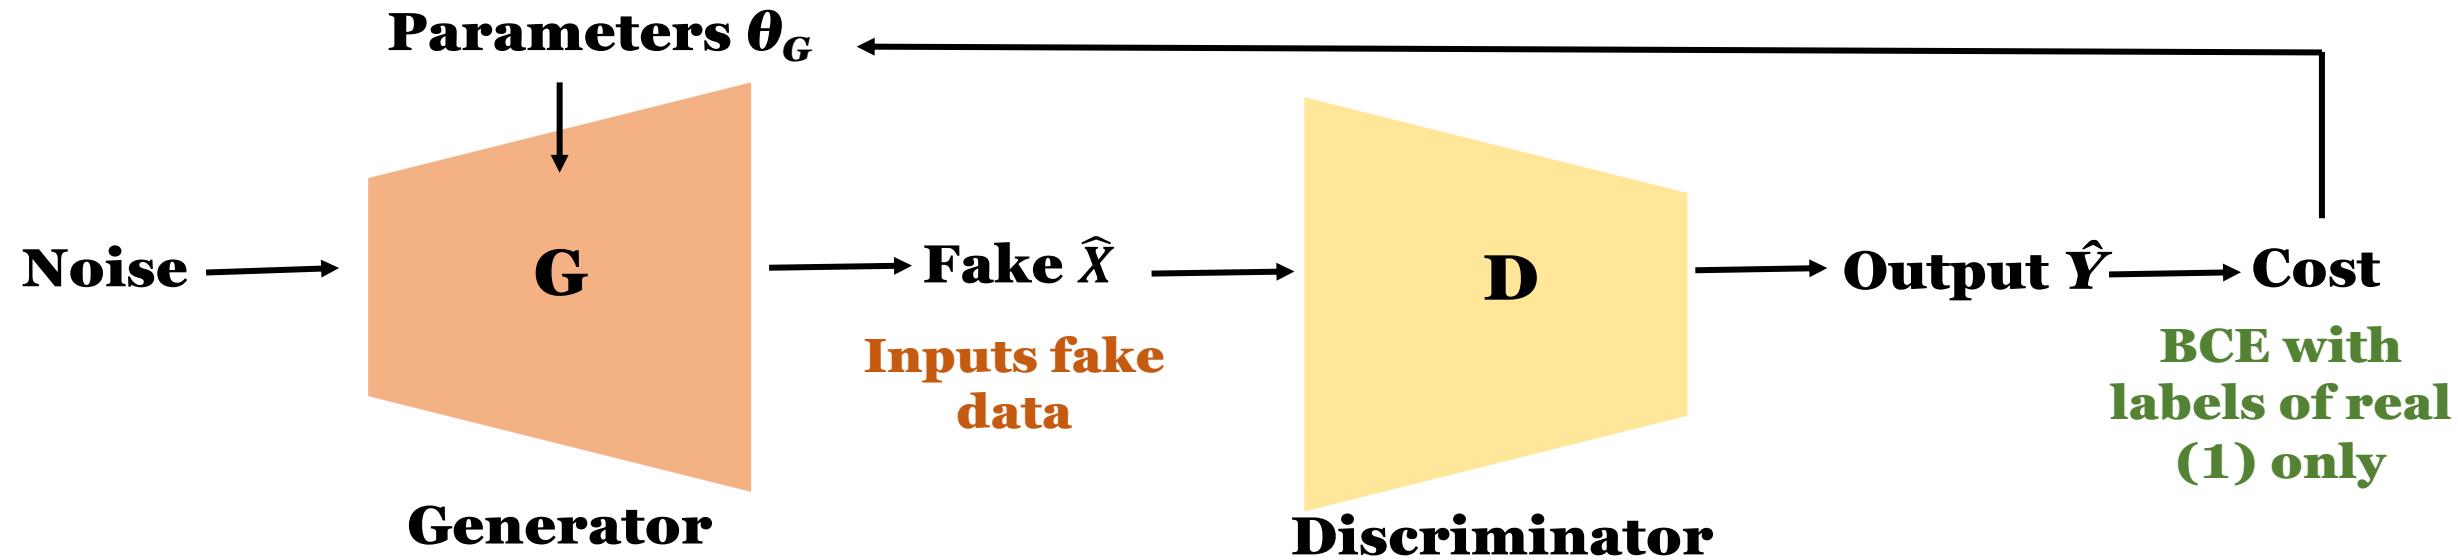
Keep generator constant



Complete GAN Model

Training Generator

Keep discriminator constant



GAN is trained in an **alternative fashion**

Do not use discriminator that is too strong!



Break Time & Quizzes

Question 1

The discriminator network and generator network influence each other solely through the data produced by the generator and the labels produced by the discriminator. When it comes to backpropagation, they are separate networks.

- True**
- False**

Break Time & Quizzes

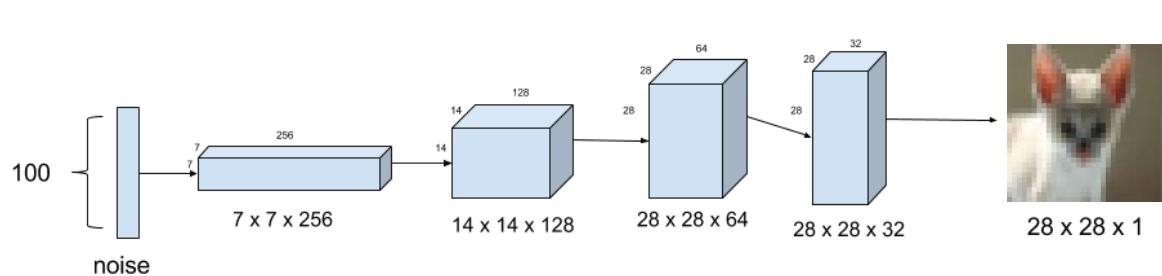
Question 2

Which is NOT TRUE about the noise vectors?

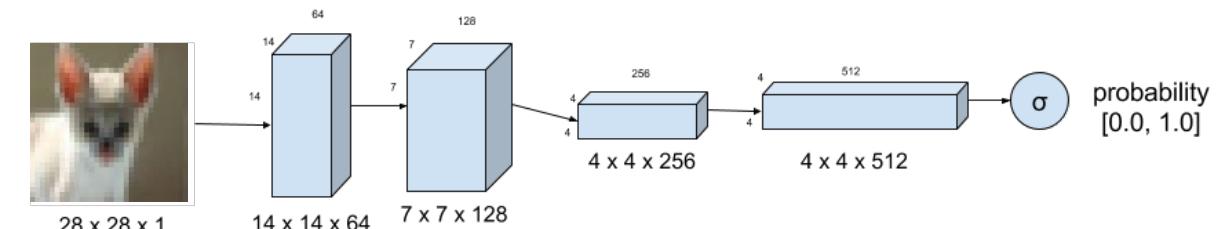
- The random noise vectors are also known as latent space vectors**
- Generator will produce different image every time we feed the same noise vector into the network**
- By introducing noise, the GAN can produce a wide variety of data, sampling from different places in the target distribution.**

DCGAN

Generator



Discriminator

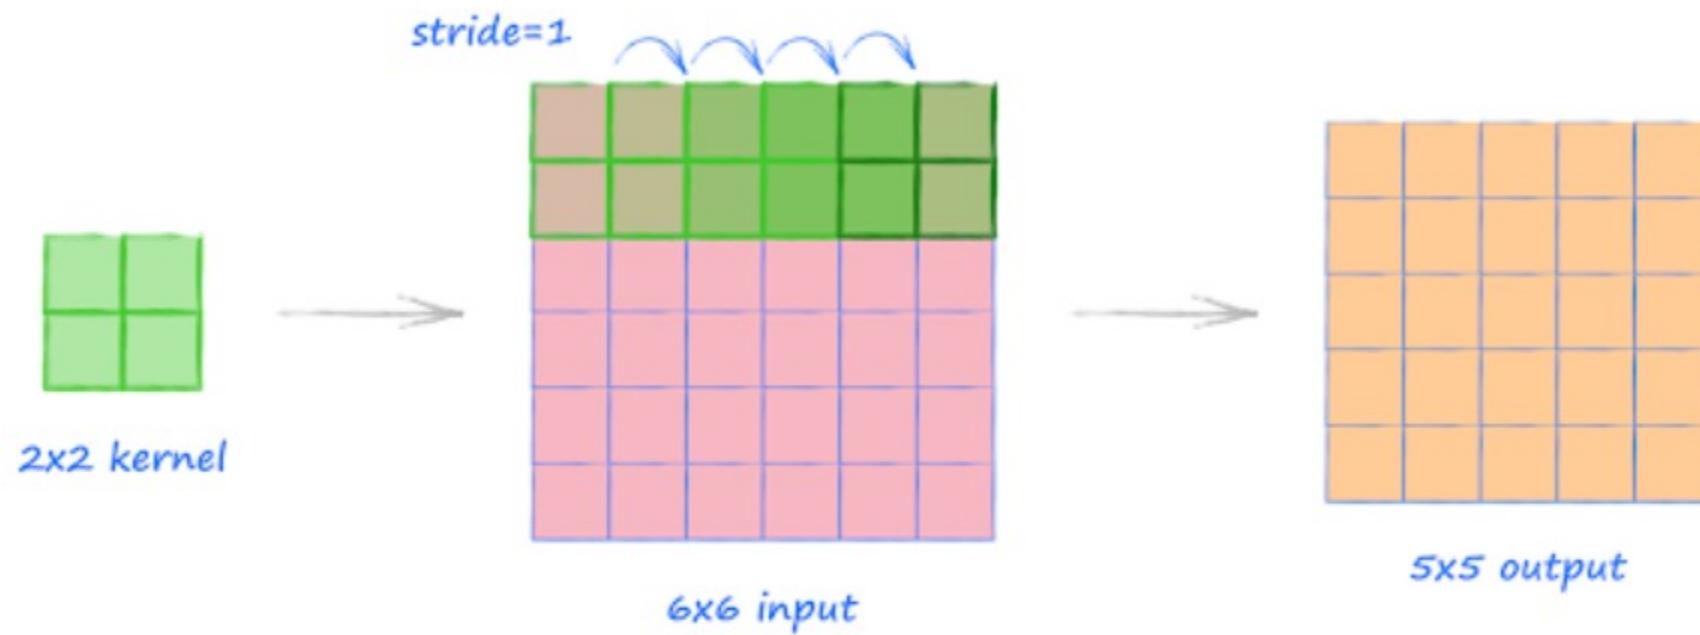


- Replace all pooling layers with convolutional stride
- Use transposed convolution for upsampling
- Use batch normalization in both the generator and the discriminator
- Remove fully connected hidden layers
- Use ReLU activation in generator for all layers except for the output, which uses Tanh
- Use LeakyReLU activation in the discriminator for all layers

Credit: <https://towardsdatascience.com/gan-by-example-using-keras-on-tensorflow-backend-1a6d515a60d0>

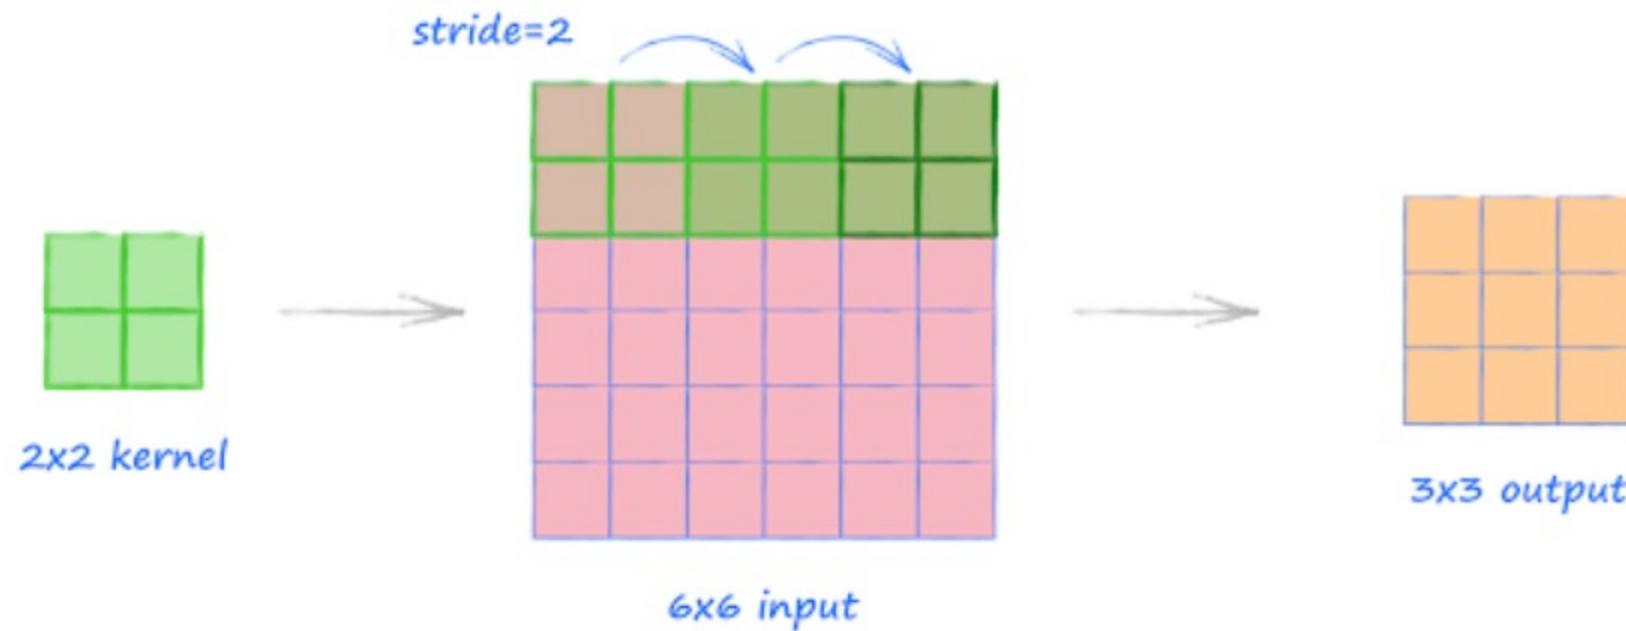
DCGAN

Discriminator: Convolution ~ downsampling



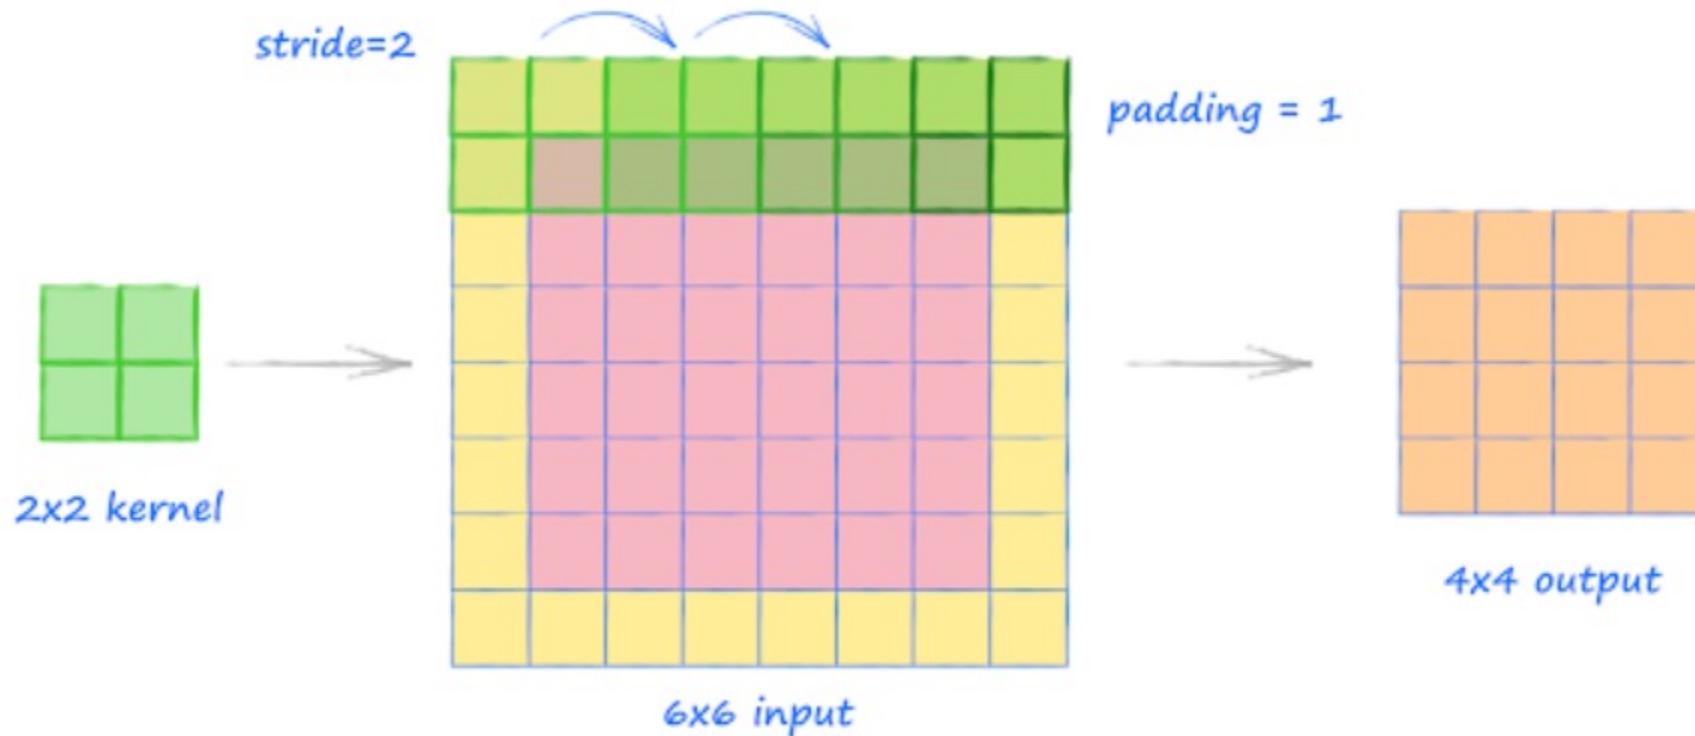
DCGAN

Discriminator: Convolution ~ downsampling



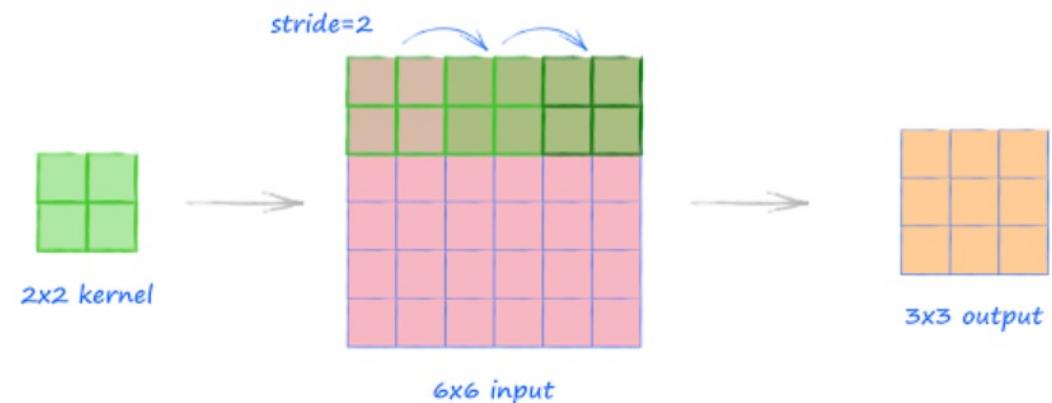
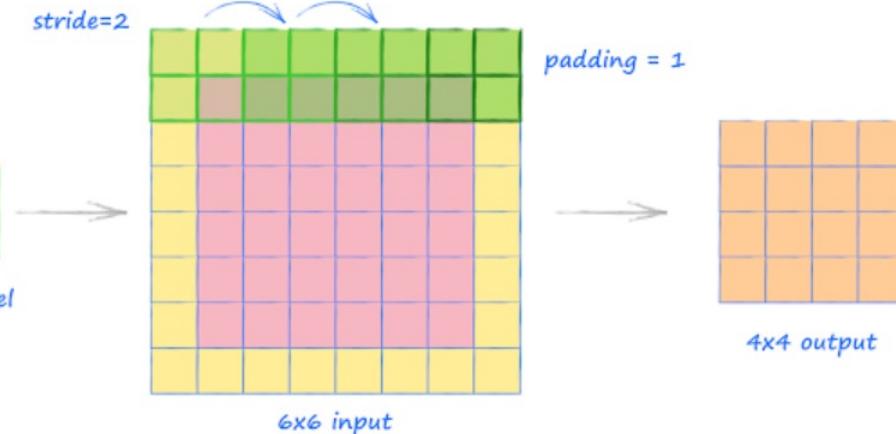
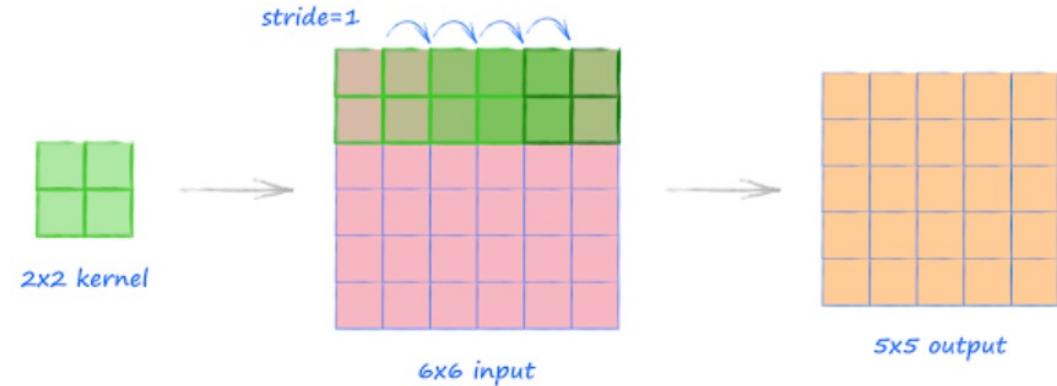
DCGAN

Discriminator: Convolution ~ downsampling



DCGAN

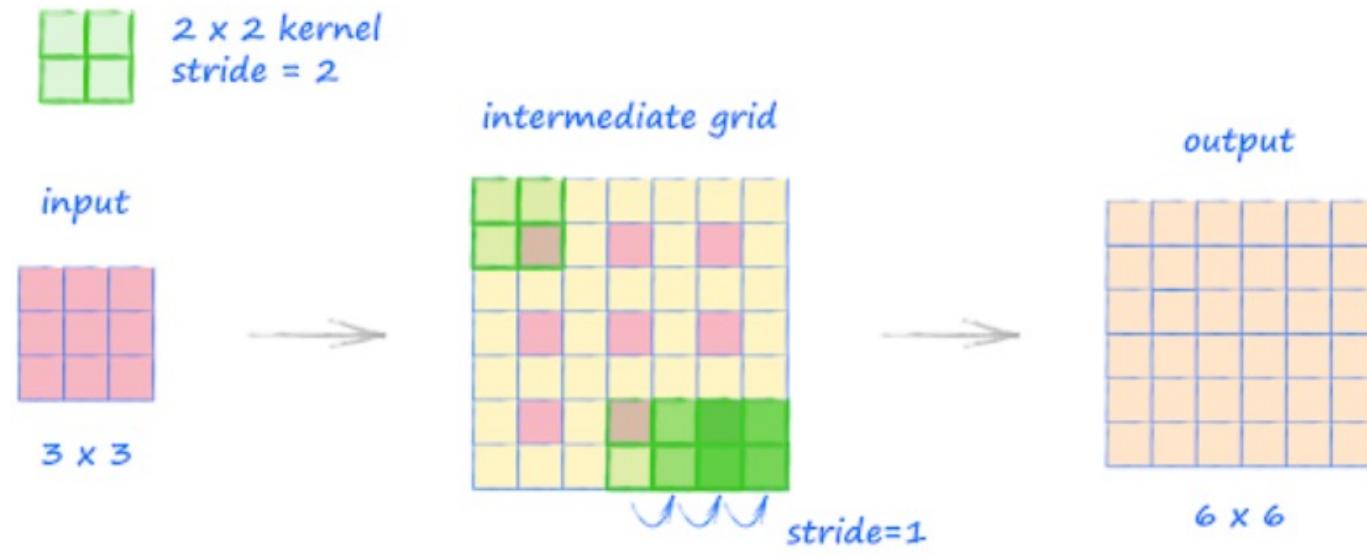
Discriminator: Convolution ~ downsampling



$$\text{output size} = \frac{\text{input size} + 2 * \text{padding} - \text{kernel size}}{\text{stride}} + 1$$

DCGAN

Generator: transposed convolution ~ upsampling

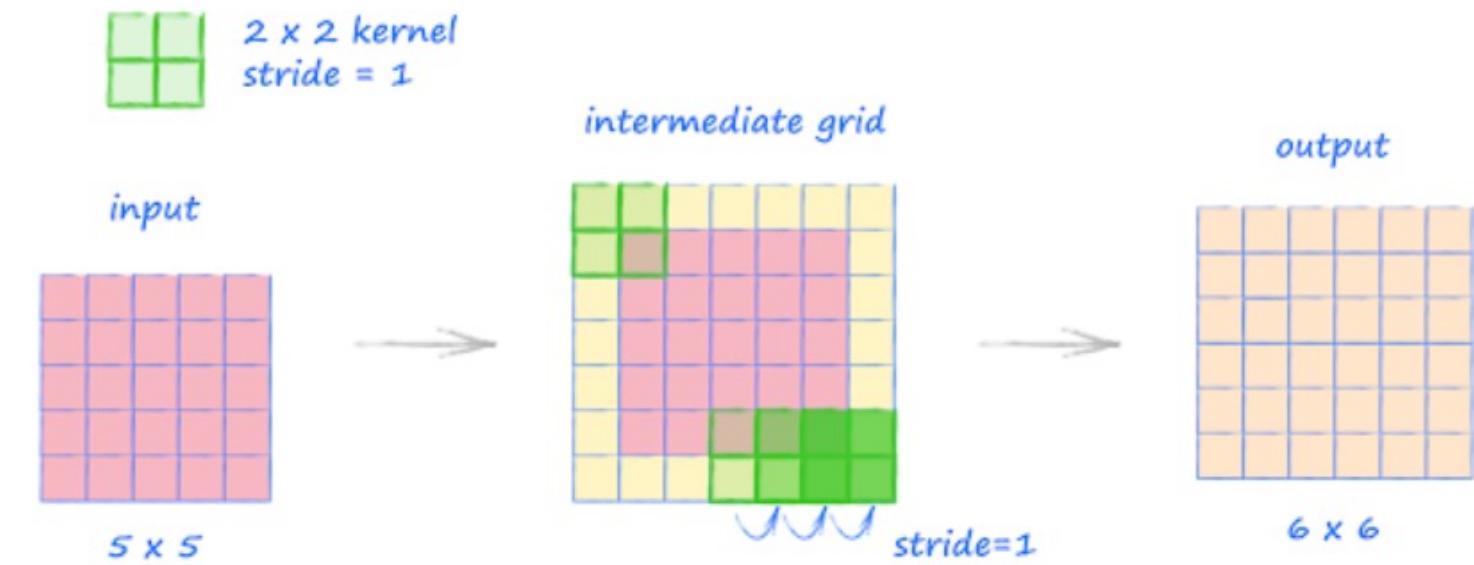


4 step process:

1. Calculate new parameters z and p'
$$z = s - 1; p' = k - p - 1$$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**

DCGAN

Generator: transposed convolution ~ upsampling

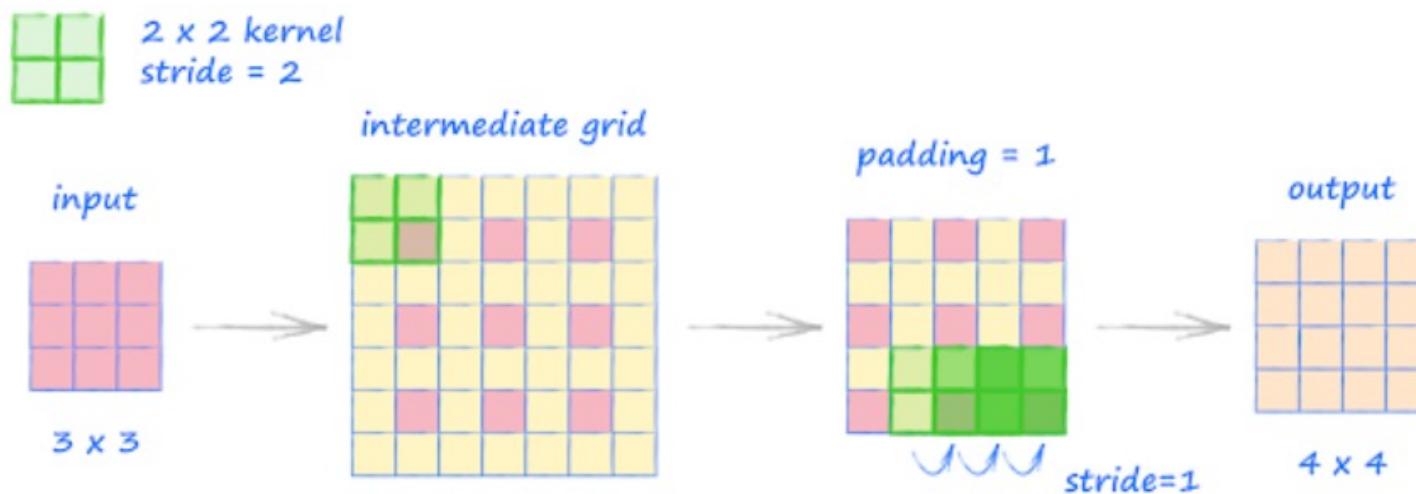


4 step process:

1. Calculate new parameters z and p'
 $z = s - 1; p' = k - p - 1$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**

DCGAN

Generator: transposed convolution ~ upsampling



4 step process:

1. Calculate new parameters z and p'
 $z = s - 1; p' = k - p - 1$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**

$$\text{output size} = (\text{input size} - 1) * \text{stride} - 2 * \text{padding} + \text{kernel size}$$

Hands-on Session

Build a DCGAN to generate CIFAR images

Exercise:

Replace ‘None’ with your own code.

THANKS FOR YOUR PARTICIPATION

We value your feedback!



<https://forms.gle/jR4Xpy2ywFtyM2wm7>

Fill in the feedback form to get the complete slides and notebook!

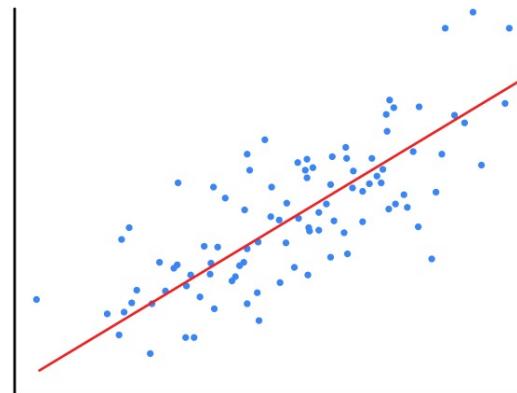
EXTRA MATERIALS

GENERATIVE ADVERSARIAL NETWORKS (GANs)

Supervised vs. Unsupervised Learning

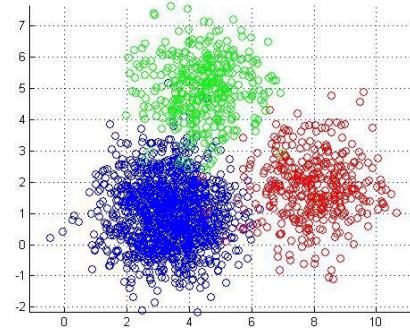
Supervised

- Data: (x, y)
 x is data, y is label
- Goal: Learn function to map $x \rightarrow y$
- Examples: Regression, classification, object detection, etc.



Unsupervised

- Data: x
 x is data, no labels
- Goal: Learn some hidden or underlying structure of the data
- Examples: Clustering, dimensionality reduction, feature learning, etc.



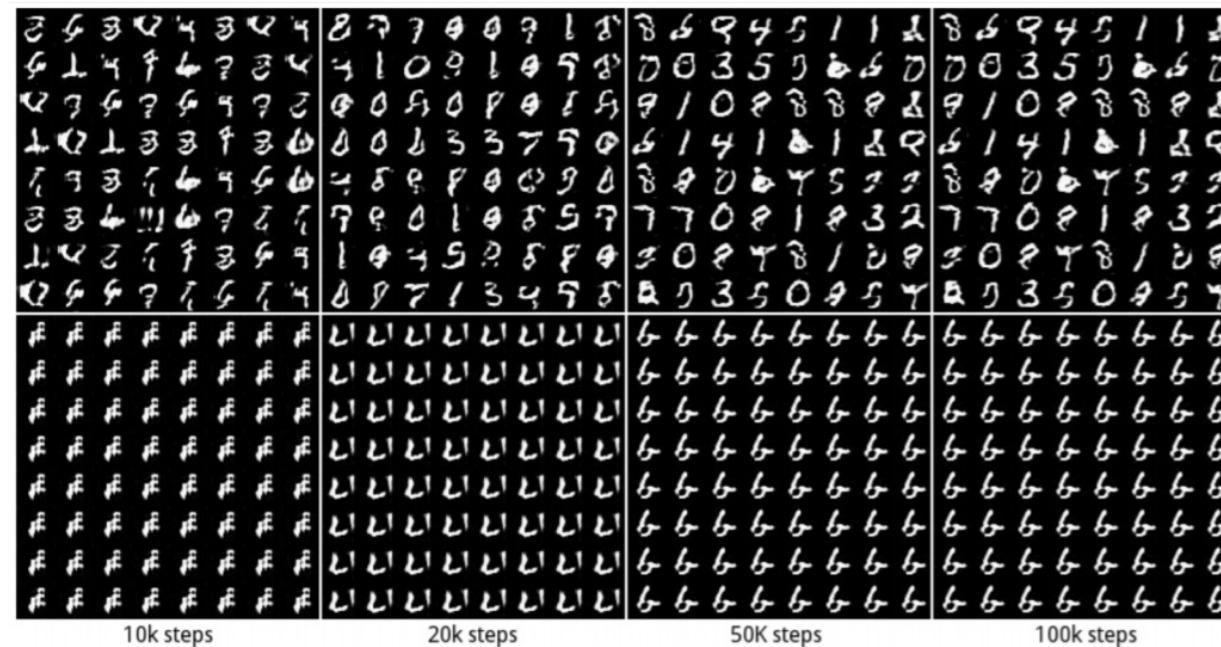
DCGAN

Batch Normalization

- Stabilize generator's learning process
- Prevent mode collapse

10 modes

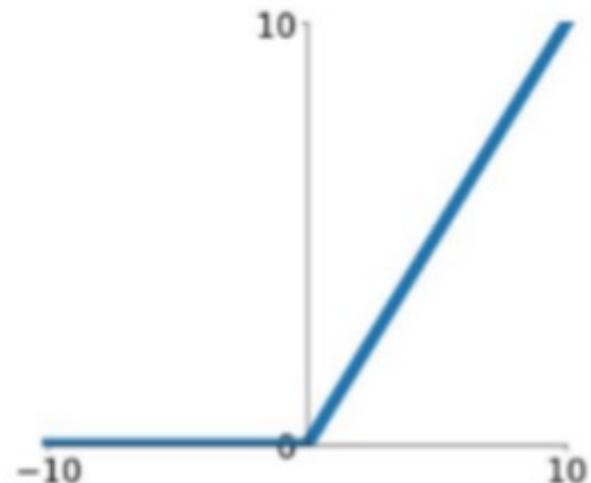
Collapse to 1 mode



DCGAN

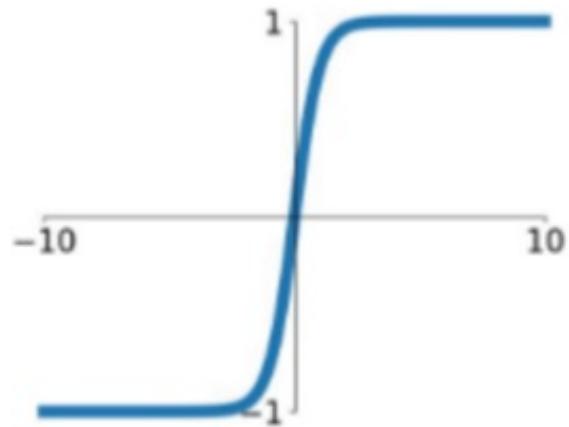
Activations:

ReLU



$$a = \max(0, z)$$

tanh

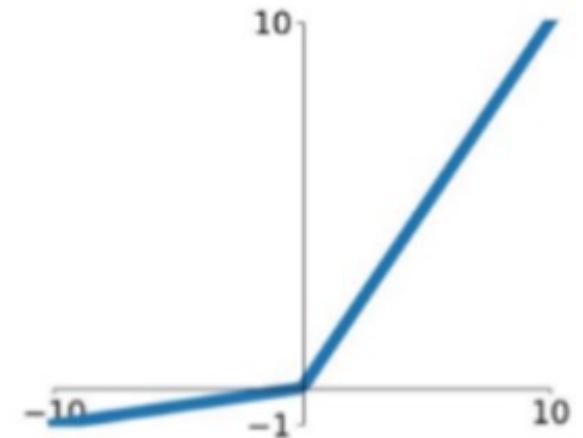


$$a = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

**Every layer of generator,
except the last layer**

Last layer of generator

Leaky ReLU



$$a = \max(0.2z, z)$$

Every layer of discriminator