


Reinforcement Learning

INSTRUCTORS: DOLPH & HOANG



MLDA
@EEE

MACHINE LEARNING AND DATA ANALYTICS



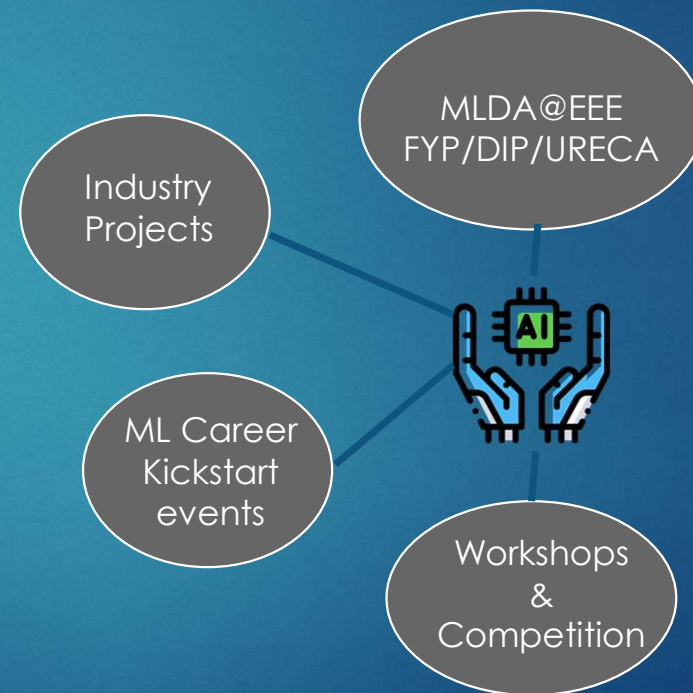
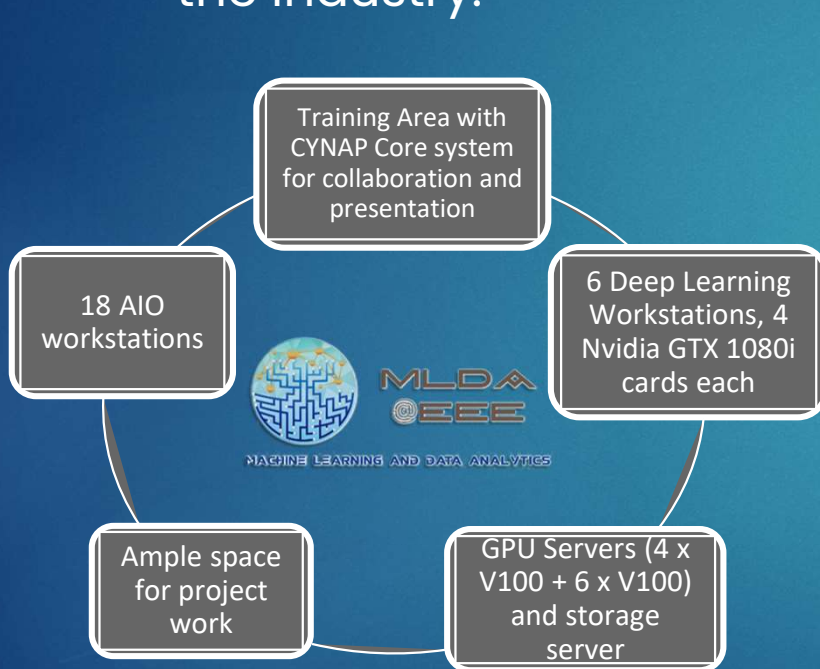
Who are we?

Dolph Xia
Tianyi
BCG Year 2

Hoang
CSC Year 2

Our Mission

Provide an integrated platform for NTU students to learn and implement Machine Learning, Data Science & AI, as well as facilitate connections with the industry.



- >1000 Trained ML practitioners
- >10 Academic Projects
- >30 Industry projects
- >5 competition
- >15 Industry Partners

Table of Contents

1

Introduction

What is RL?

2

Markov Decision Process

Key terminologies

3

Algorithms

Classic techniques

4

Hands-on

Tinkering time!



1

Introduction

What is reinforcement
learning (RL)?



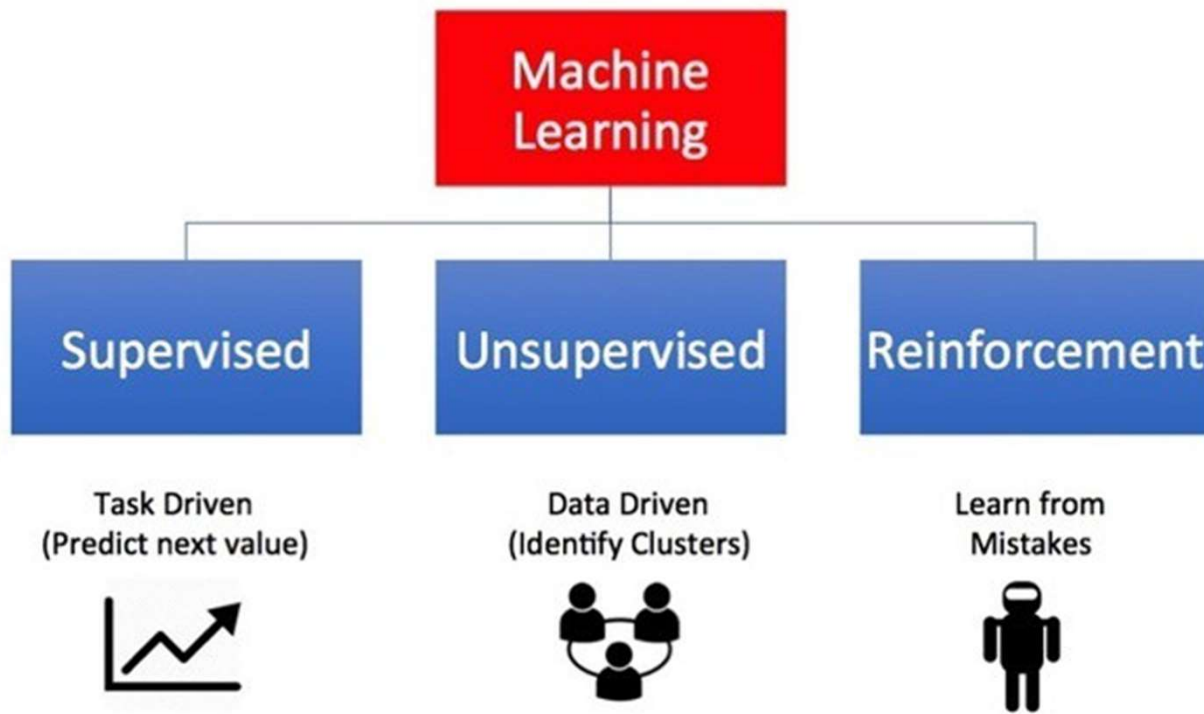
How you learn to play a new game?



- ▶ None/ minimal instructions
- ▶ Just try out different combinations and see how they work out

What is Reinforcement Learning?

Types of Machine Learning



The agent learns by interacting with the environment.

Goal: **maximize the total reward**

- **Trial and error** search for optimal actions
- Actions affect immediate and subsequent rewards. **Feedback can be delayed**

RL Application: Game



Clear goals, sequential decision making

“This indicates that reinforcement learning can yield long-term planning with large but achievable scale — without fundamental advances, contrary to our own expectations upon starting the project.”



RL Application: autonomous driving




RL Application: Automated system



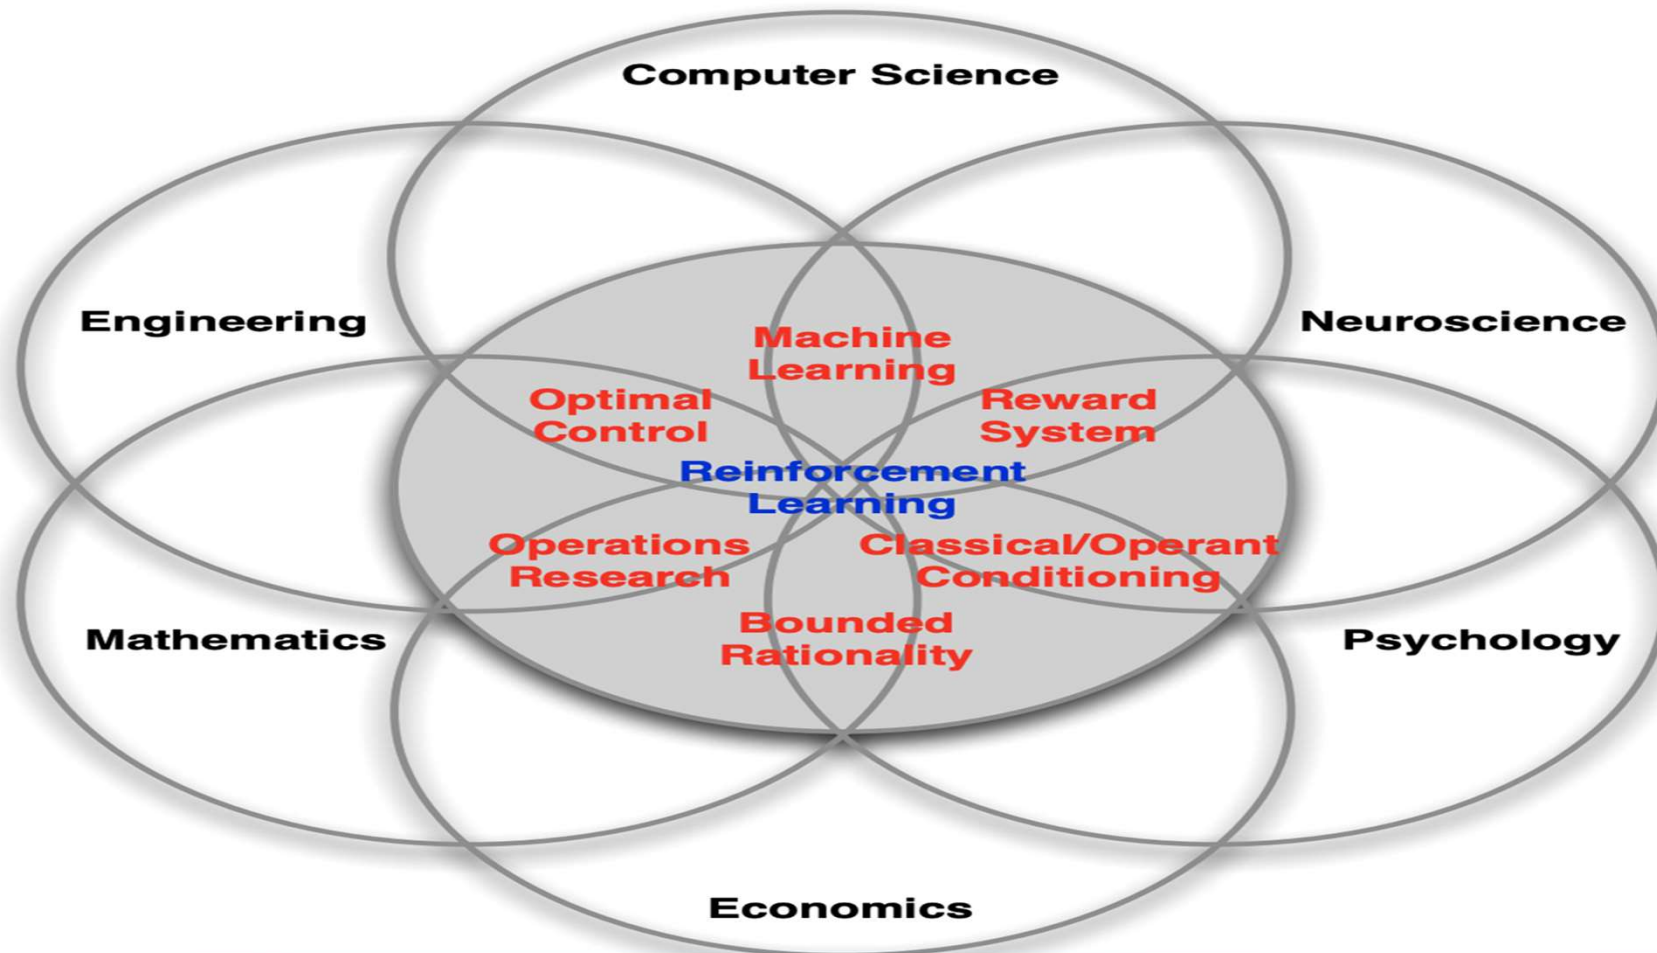
**DeepMind AI Reduces
Google Data Centre
Cooling Bill by 40%**



RL Application: robot



Many Facets of Reinforcement Learning





2

Markov Decision Process

Key Terminologies



How to model the problem?



How to decide what to do?

- The current state
- Available options
- Effect of actions on the environment

Markov Decision Process



$\langle S, A, P, R, \gamma \rangle$

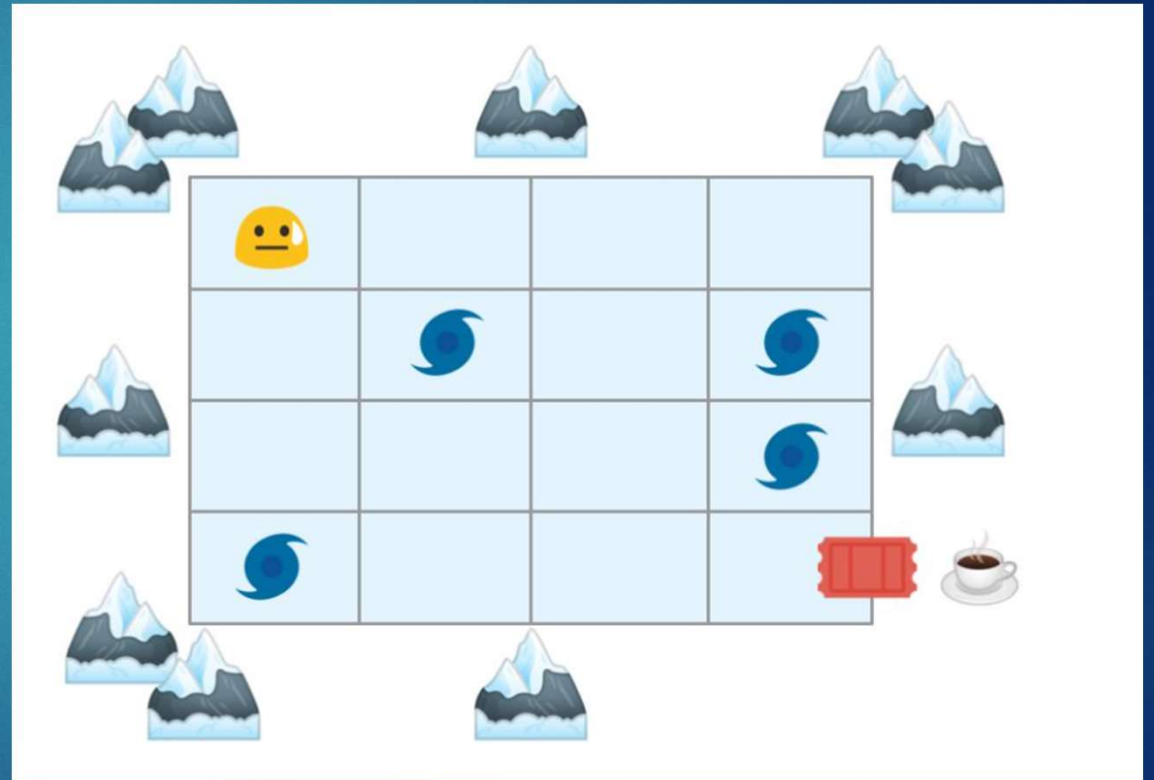
S: **State**

A: **Action**

P: **Transition**: how the env changes

R: **Reward**

γ : **discount factor**: make total
reward converge



Markov Decision Process



Policy π : The strategy of the agent

The action for each state

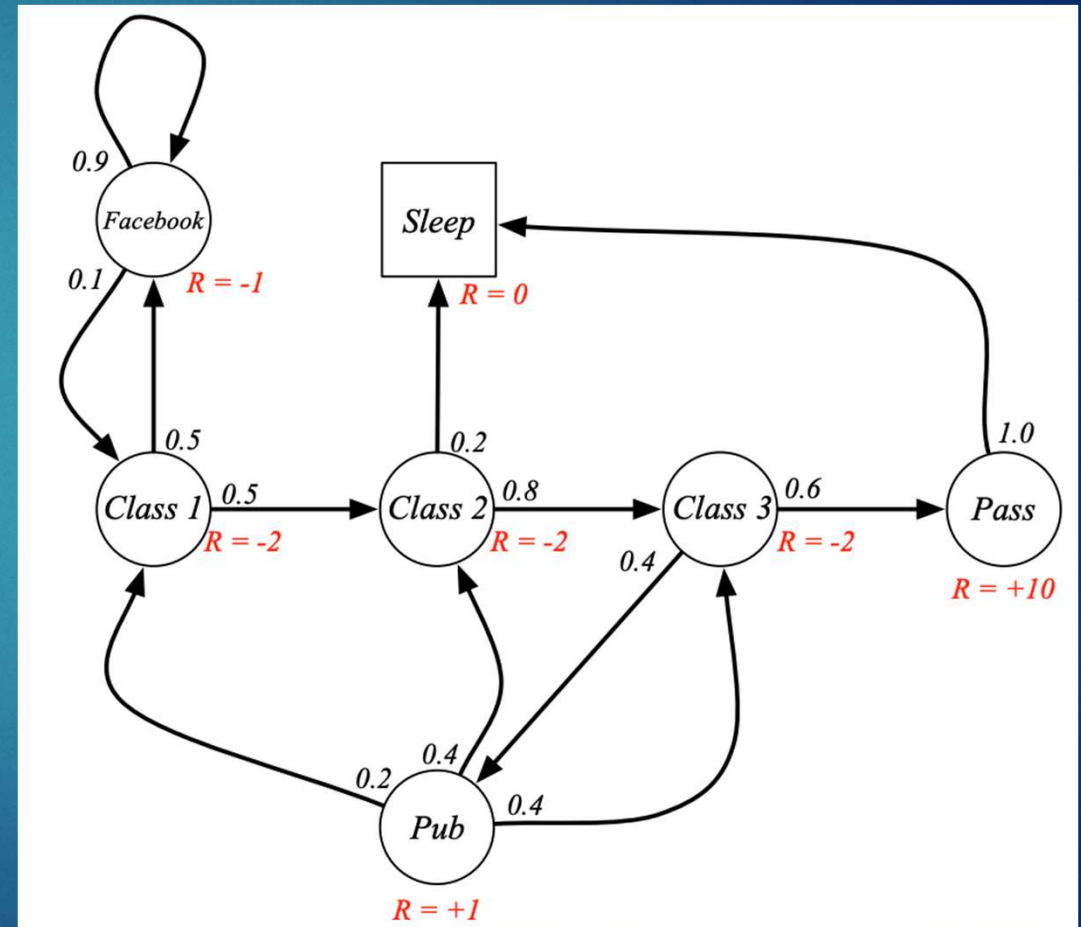
Example 1:

“always attend lectures”

Example 2:

“follow the probability shown in the picture”

Return: total discounted reward till terminal state




Markov Decision Process



State value function: $v_{\pi}(s)$

The expected return from state s ,
under policy π

Example :
Value of 'Pass' is 10.



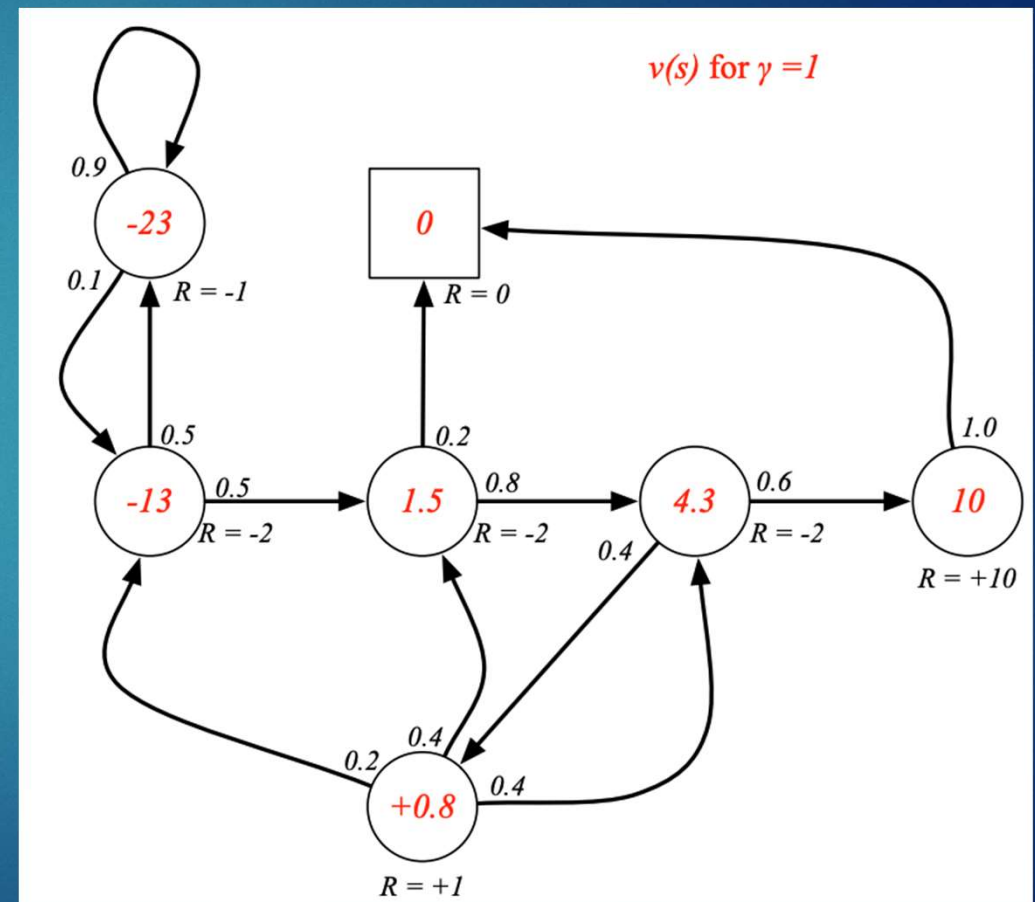
Markov Decision Process



Action value function: $q_{\pi}(s, a)$

The expected return from state s ,
taking action a , under policy π

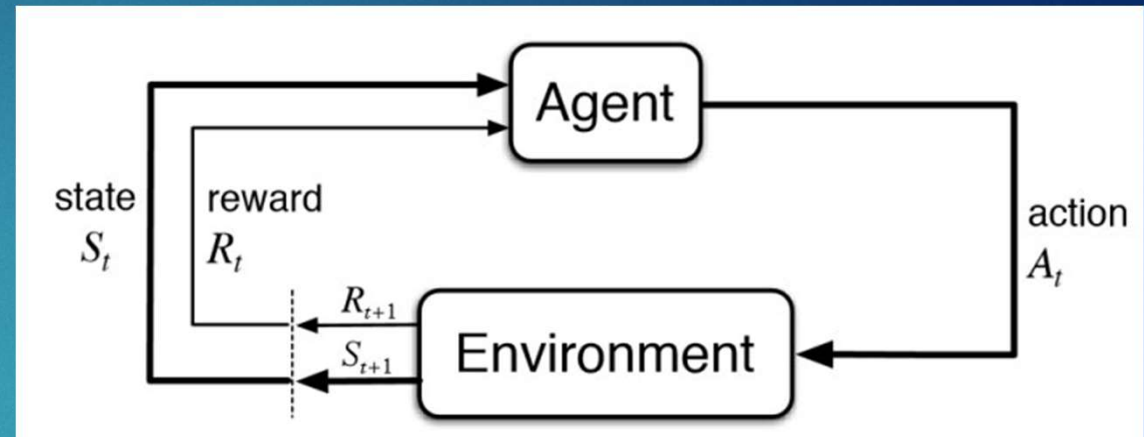
Example :
Value of 'Stay in Facebook state' is -23.



Markov Decision Process



Sequential decision making



Future states are **only dependent on the present** state


$\langle S, A, P, R, \gamma \rangle$: state, action, transition, reward, discounting factor

$$\mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

An example

State: index of box (0 for start state) **Action:** 0 left 1 down 2 right 3 up

Start
state



Return: total discounted future rewards

Terminal
states

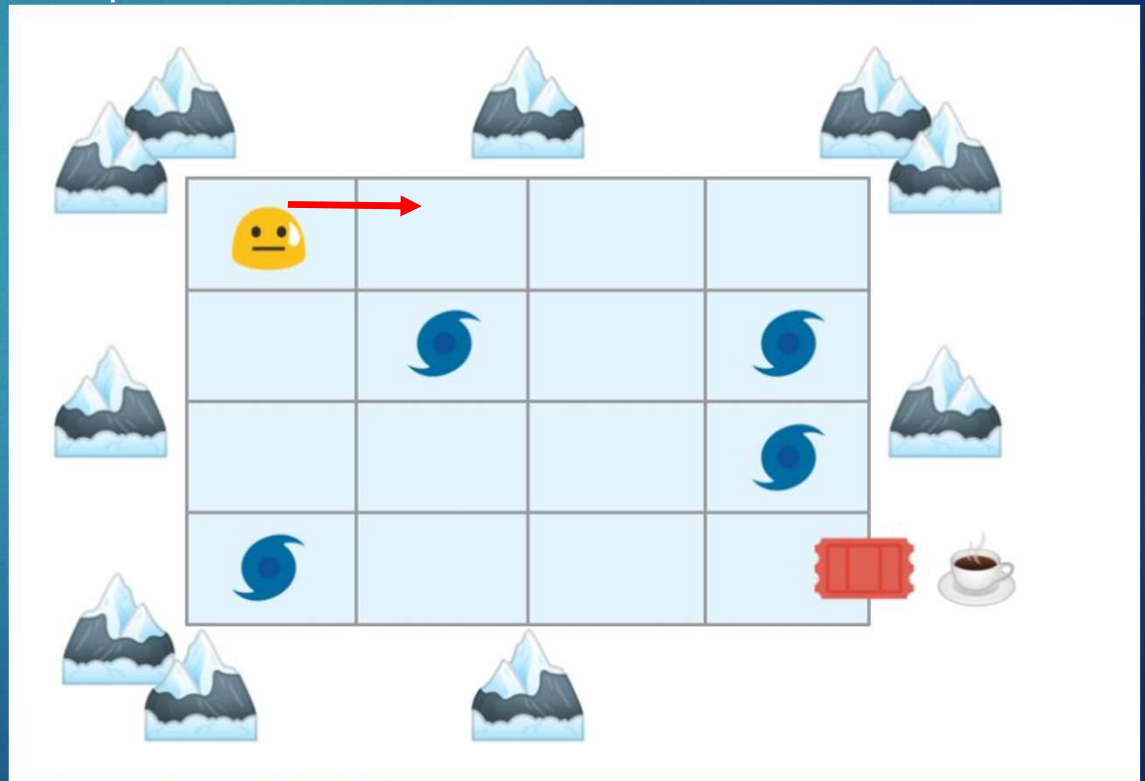
Policy

A mapping from states to ~~actions~~ **More general:** “probabilities of selecting each possible action”

Example 1:
“always attend lectures”

Example 2:
“always move right”

Example 3:
“roll a 6-sided dice
and move right if get 6”



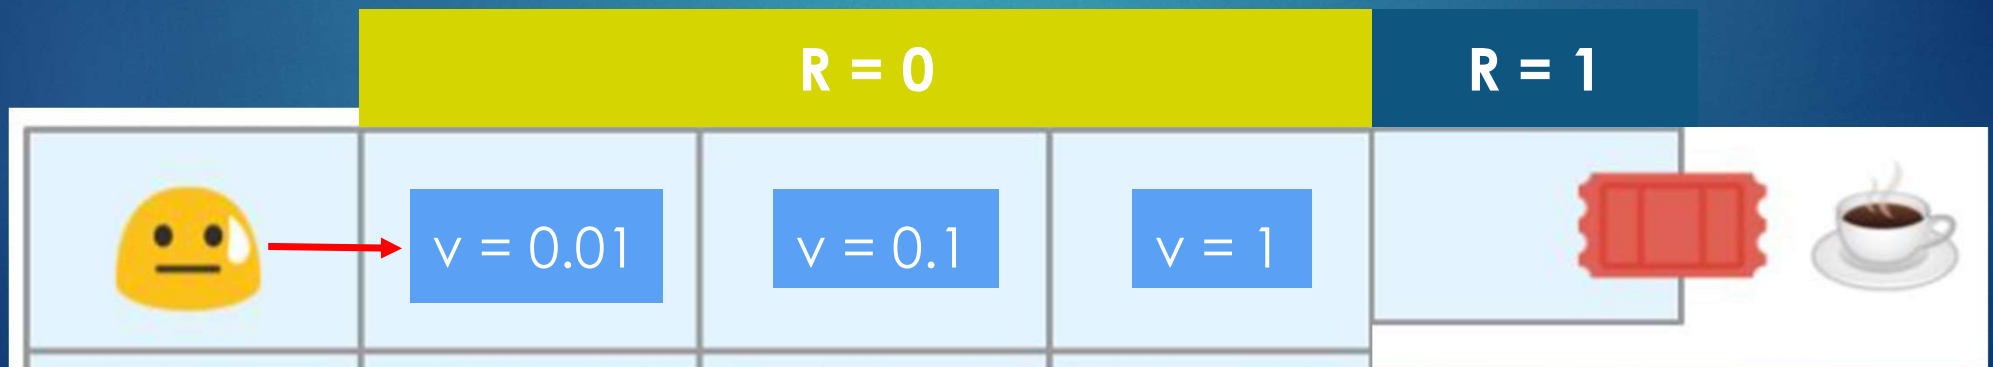
State-Value function

Expected return from **state**, following a **particular policy**.

Value of state s for policy π :

$$v_{\pi}(s)$$

Example: policy “always move right”, with discount factor = 0.1



Action-Value function

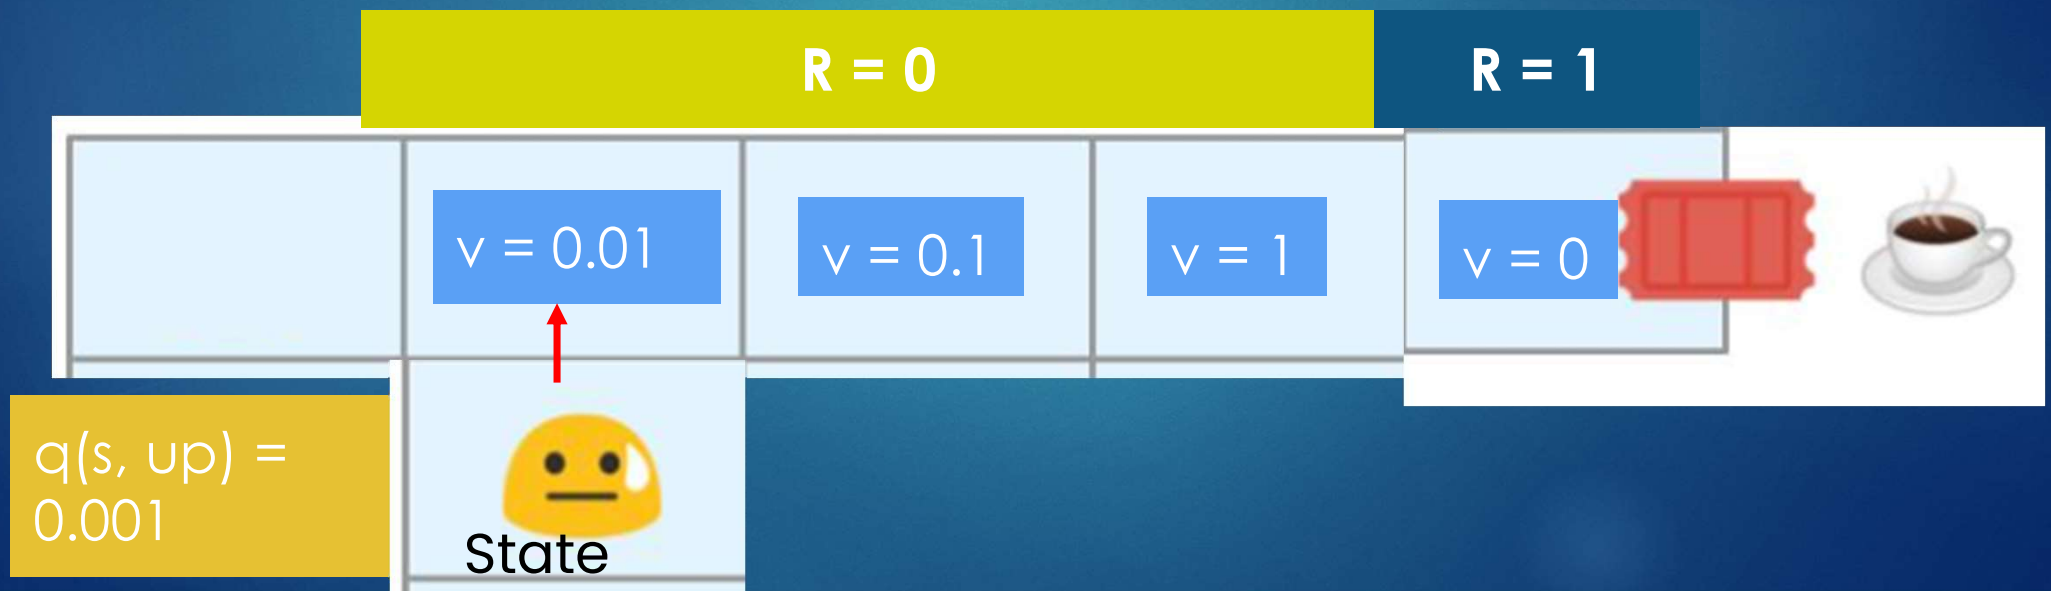
Expected return from **state**, by **taking an action**, thereafter **following a particular policy**.

Action-Value of state s , action a , policy π :

$$q_{\pi}(s, a)$$

Note that action may be arbitrary and doesn't have to be in line with the policy

Example: policy "always move right"



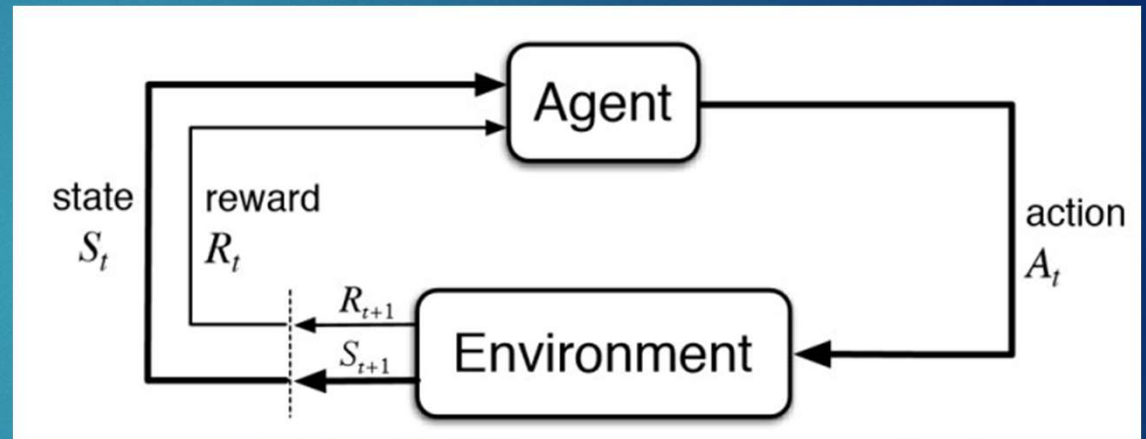
MDP summary

Return: discounted total reward

State value: expected return from state s

Action value: expected return by action a

Policy: action distribution for each state




Bellman Equation

To calculate expected total discounted reward by a recursive function.

$$v(s) = \mathbb{E}[G_t \mid S_t = s]$$

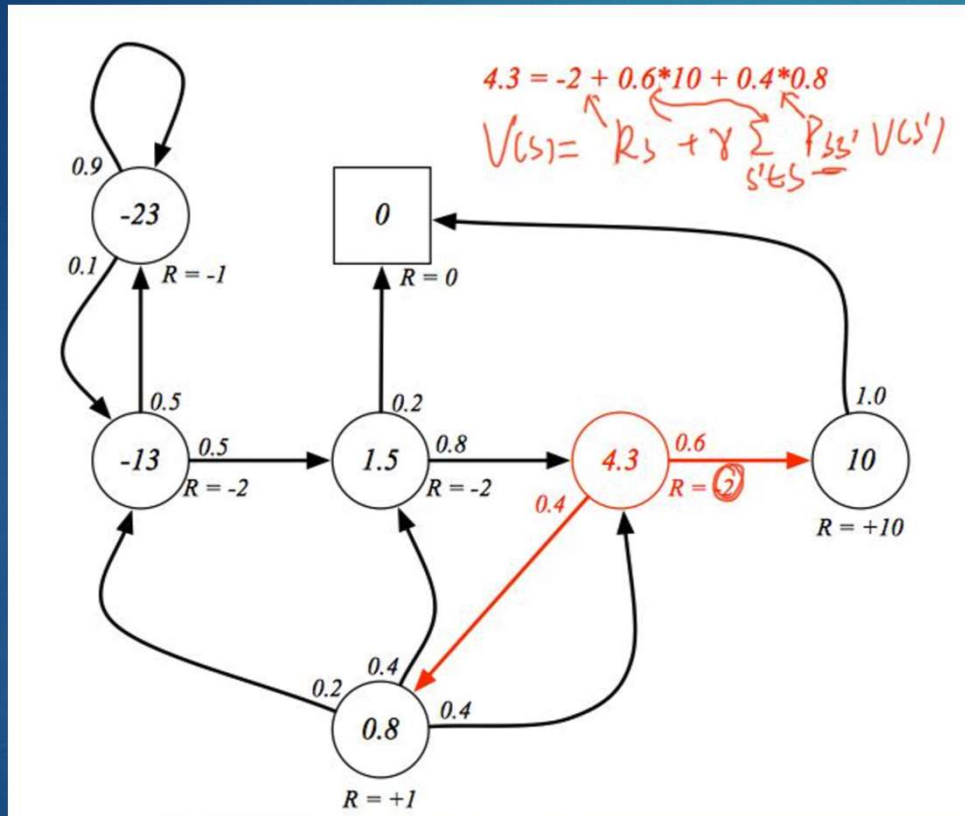
$$= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s]$$

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$



The state-value function and action-value function can be decomposed into immediate reward plus discounted value of successor state

Bellman Equation Example



Exploitation and exploration


Motivation: find global optimal solution

Greedy: always select best known (may miss global optimum)

Example: select a restaurant

Go to favourite restaurant or try a new restaurant

epsilon-greedy: with a small ratio, choose an action randomly rather than the currently optimal action





3

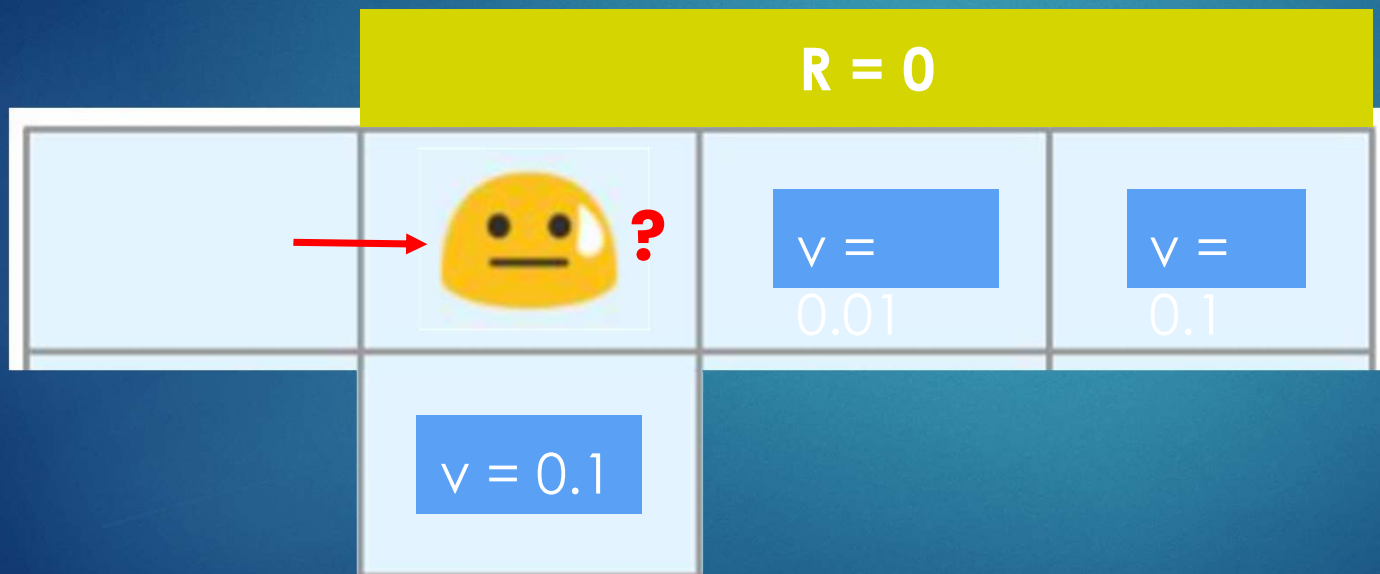
Algorithms

Classic techniques



How to improve policy?

Given: value function of a policy (always move right)



Policy Improvement Theorem

Eureka! 

If we “improve” action at one state s , with respect to v and keep everything else same, we get a better policy than π

Smaller Eureka:
The actual theorem is a bit more complicated.

Generalized Policy Iteration

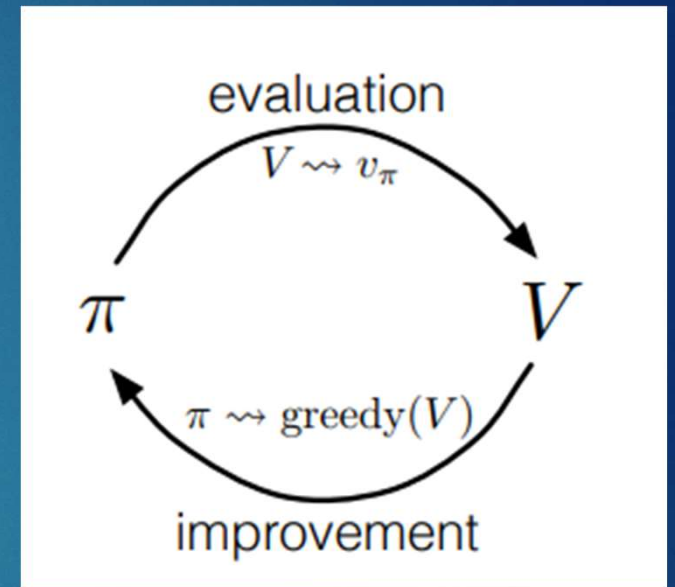
1. initialize policy

Loop:

2. Evaluate value function of policy

3. Improve policy by greedifying (to certain extent) with value function

Loop until policy is greedy with respect to its own value function



Dynamic Programming

- 1) Assume a **complete accurate model** of the environment


If we choose to move right...



We know exactly **where** the gumdrop may end up,
with **what probability**

Dynamic Programming


- 1) Assume a **complete accurate model** of the environment
- 2) Update state values based on **estimated** values of other states: *bootstrapping*



$$v(s) = 50\% * (0 + v(s_1)) + 50\% * (1 + v(s_2))$$
$$= 0.55$$

Dynamic Programming

Visualizing
bootstrapping



Dynamic Programming

Notable examples

- 1) Policy Iteration
- 2) Value Iteration

Policy Iteration - How?

1. initialize policy


... by bootstrapping **till**
converge

Loop:

2. Evaluate value function of policy


3. Improve policy by greedifying (to certain extent) with value function

Loop until policy is greedy with respect to its own value function




Policy Iteration: An Illustration

Iteration 1

Policy		
→	→	→
→	→	
→	→	→

Expected Value		
0	0	0
0.1	1	0
0	0	0

Iteration 2

↓	↓	→
→	→	
↑	↑	→

0.01	0.1	0
0.1	1	0
0.01	0.1	0

Value iteration: How?

Initialise state-value for each state

Loop until convergence:

For each state, update state-value as the highest action-value

Policy is greedy with respect to the converged state-value function

Value Iteration - An illustration

Initialise value function

0	0	0
0	0	0
0	0	0

Iteration 1

0	0	1
0	1	0
0	0	1

Iteration 2


0	0.1	1
0.1	1	0
0	0.1	1

No more change after
Iteration 3

↓	↓	↓
→	→	☕
↑	↑	↑

Limitations in Dynamic Programming

- ↩ Can't solve unknown environment




Monte Carlo


Differs from DP in how it evaluates value function

Learn value functions from experience

- Generate numerous experiences by following the policy
- Value of a state = average of the returns



Number of points	Approximate area (should be 3.14159...)
10	2.4
100	3.0
1000	3.152
10000	3.14
10000	3.14896
100000	3.14112



Q-Learning


Q learning is a value-based method of supplying information to inform which action an agent should take.

Directly approximate action-value function of **optimal policy**

Over-simplified example:

$q(\text{today}, \text{studyRL}) \leftarrow \text{value of reward from this workshop} +$
+ value for studying “greedily” forever

$q(\text{today}, \text{play}) \leftarrow \text{value of reward from playing} +$
+ value for studying “greedily” forever



Q-Learning

Initialize all q values, store in table

Bellman equation for update

Current value

Reward

Maximum reward that can be obtained from state

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

Policy Gradient

Optimal action may not be deterministic.

Optimizing parametrized policies with respect to the expected return by gradient descent.

No value function required

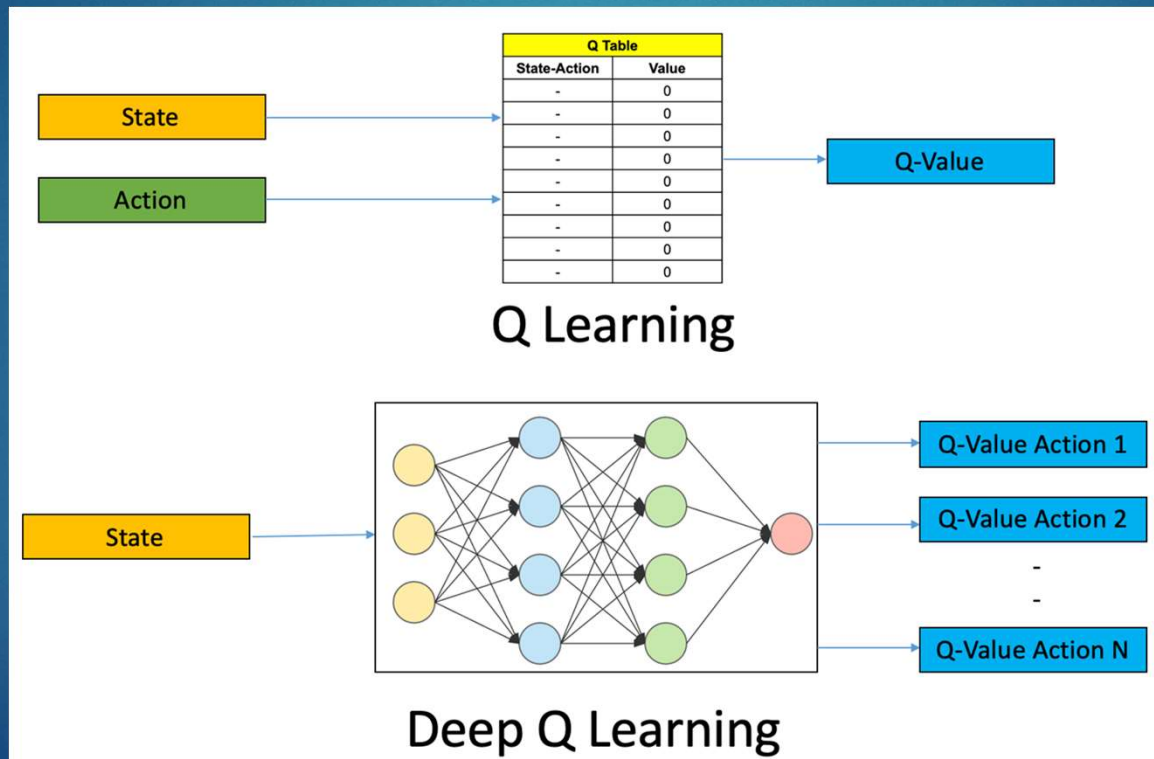
A deterministic policy can be easily exploited

Uniform random policy



100

Environments and actions can be continuous



Break & some questions for you!





4

Hands-on

Link will be sent in Zoom chat



Question 1:

What is NOT a characteristic of reinforcement learning:

Option 1: It is a sequential decision making procedure.

Option 2: It is a branch of machine learning.

Option 3: Reward signals and feedbacks are always immediate.

Option 4: Time plays a crucial role in reinforcement learning.

Question 2:

Which of the following is NOT true on Markov Decision Process(MDP)?

Option 1: Agent is the learner and the decision maker.

Option 2: At each time step the agent takes an action.

Option 3: At each time step the environment generates a reward signal.

Option 4: All the past history states and actions are required to determine the next state.

Question 3:

Suppose discount factor $\gamma=0.8$, we observe the following sequence of rewards:

$R_1 = -3, R_2 = 5, R_3=2, R_4 = 7$, and $R_5 = 1$, with $T=5$. What is the return G_0 ? Hint: Work Backwards and recall that $G_t = R_{t+1} + \gamma * G_{t+1}$.

Option 1: 5.27

Option 2: 6.27

Option 3: 7.27

Option 4: 8.27

Question 3:

Which of the algorithms requires that the agent has complete knowledge of the environment?

Option 1: Value iteration.

Option 2: Monte Carlo.

Option 3: Q learning.

Option 4: Policy gradient.


Any questions for us?



Moving forward

- ⤴ Going broad: learn other fundamental concepts and techniques
- ⤴ Going deep: Deep Reinforcement Learning
- ⤴ Going practical: train the models on some real games!

Where are the **links** to do all this? => They are in the feedback form 😊



Thanks!

Your feedback is extremely valuable to **reinforce our learning** as instructors, to create **better workshops for you** in the future! The **link to the additional materials** is available after submitting the feedback form.

<https://forms.gle/GR4YnphhvEbmXoXv9>

Credits

- ↳ <https://docs.paperspace.com/machine-learning/wiki/supervised-unsupervised-and-reinforcement-learning>
- ↳ https://raw.githubusercontent.com/FrancescoSaverioZuppichini/Value-Iteration-Network/master/core/gridworld_28x28/animation.gif
- ↳ Sutton, R. S., Barto, A. G. (2018). Reinforcement Learning: An Introduction. MIT Press.
- ↳ <https://www.davidsilver.uk/teaching/>
- ↳ Mykel J. Kochenderfer (2021). Algorithms for Decision Making .MIT Press.
- ↳ <https://dreager1.files.wordpress.com/2010/05/mario-vs-browser.jpg>
- ↳ <https://i.insider.com/560ebbe7dd0895325c8b458e?width=1100&format=jpeg&auto=webp>
- ↳ <https://www.analyticsvidhya.com/blog/2019/04/introduction-deep-q-learning-python/>